

# Sockets Introduction

- Socket address structures
- Value-result arguments
- Byte ordering and manipulation functions
- Address conversion functions: *inet\_aton*,  
*inet\_addr*, *inet\_ntoa*, *inet\_pton*, *inet\_ntop*,  
*sock\_ntop*
- Stream socket I/O functions: *readn*, *writen*,  
*readline*
- File descriptor testing function: *isfdtype*

# Socket Address Structures: IPv4, Generic, IPv6

```
struct in_addr {  
    in_addr_t s_addr; /* 32-bit IPv4 address, network byte order */};  
  
struct sockaddr_in {  
    uint8_t sin_len; /* length of structure */  
    sa_family_t sin_family; /* AF_INET */  
    in_port_t sin_port; /* 16-bit port#, network byte order */  
    struct in_addr sin_addr; /* 32-bit IPv4 address, network byte order */  
    char sin_zero[8]; /* unused */};  
  
struct sockaddr { /* only used to cast pointers */  
    uint8_t sa_len;  
    sa_family_t sa_family; /* address family: AF_xxx value */  
    char sa_data[14]; /* protocol-specific address */};  
  
struct in6_addr {  
    uint8_t s6_addr[16]; /* 128-bit IPv6 address, network byte order */};  
  
struct sockaddr_in6 {  
    uint8_t sin6_len; /* length of this struct [24] */  
    sa_family_t sin6_family; /* AF_INET6 */  
    in_port_t sin6_port; /* port#, network byte order */  
    uint32_t sin6_flowinfo; /* flow label and priority */  
    struct in6_addr sin6_addr; /* IPv6 address, network byte order */};
```

# Datatypes Required by Posix.1g



Datatype	Description	Header
int8_t	signed 8-bit integer	<sys/types.h>
uint8_t	unsigned 8-bit integer	<sys/types.h>
int16_t	signed 16-bit integer	<sys/types.h>
uint16_t	unsigned 16-bit integer	<sys/types.h>
int32_t	signed 32-bit integer	<sys/types.h>
sa_family_t	address family of socket addr struct	<sys/types.h>
socklen_t	length of socket addr struct, uint32_t	<sys/types.h>
in_addr_t	IPv4 address, normally uint32_t	<sys/types.h>
in_port_t	TCP or UDP port, normally uint16_t	<sys/types.h>

## Figure 3.3 The generic socket address structure: `sockaddr`.

```
struct sockaddr {
    uint8_t      sa_len;
    sa_family_t   sa_family;    /* address family: AF_xxx value */
    char         sa_data[14];   /* protocol-specific address */
};
```



```
int bind(int, struct sockaddr *, socklen_t);
```

This requires that any calls to these functions must cast the pointer to the protocol-specific socket address structure to be a pointer to a generic socket address structure. For example,

```
struct sockaddr_in     serv;      /* IPv4 socket address structure */
/* fill in serv{} */

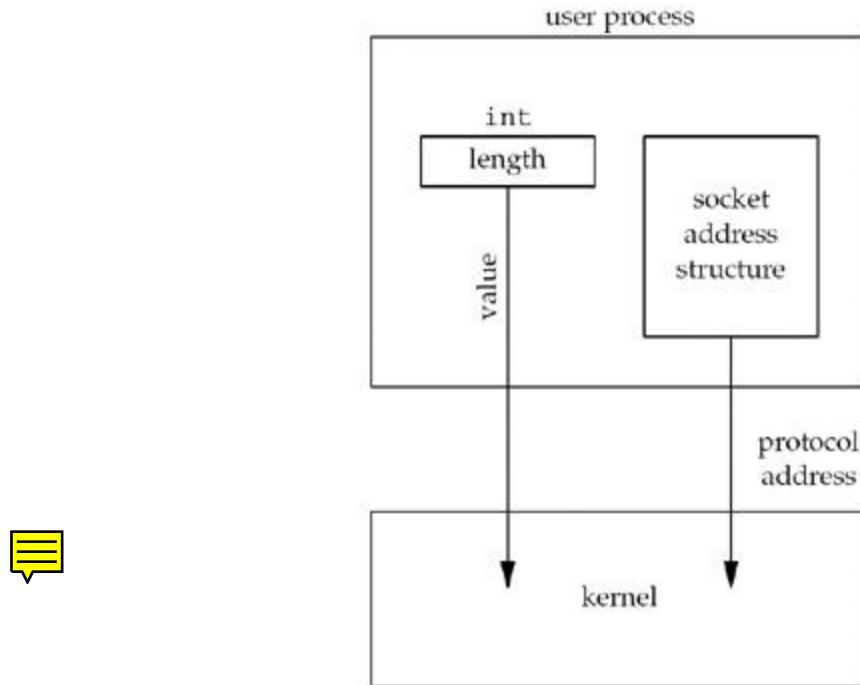
bind(sockfd, (struct sockaddr *) &serv, sizeof(serv));
```

# Comparison of Various Socket Address Structures

IPv4 <code>socketaddr_in{}</code>	IPv6 <code>socketaddr_in6{}</code>	UNIX <code>socketaddr_un{}</code>	Datalink <code>socketaddr_dl{}</code>
length AF_INET 16-bit port# 32-bit IP address (unused) fixed length (16 bytes)	length AF_INET6 16-bit port# 32-bit flow label 128-bit IPv6 address fixed length (24 bytes)	length AF_LOCAL pathname (up to 104 bytes)	length AF_LINK interface index type name len addr len sel len interface name and link-layer address variable length
			variable length

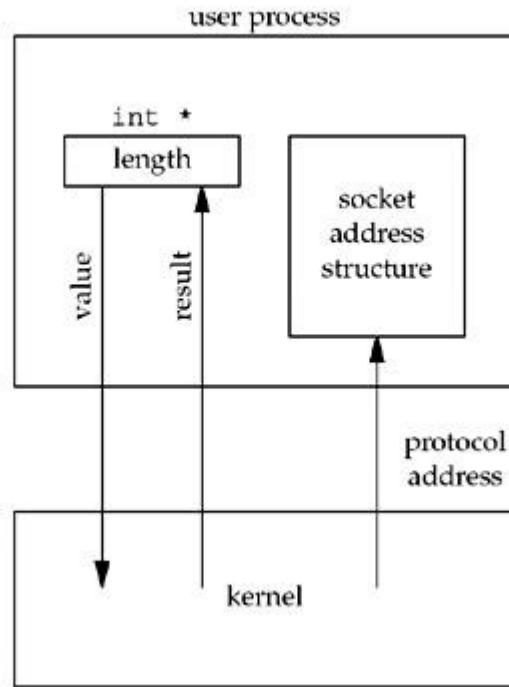
# Value-Result Argument

Figure 3.7. Socket address structure passed from process to kernel.



```
struct sockaddr_in serv;  
/* fill in serv{} */  
connect(sockfd, (SA *) &serv, sizeof(serv));
```

Figure 3.8. Socket address structure passed from kernel to process.



```
struct sockaddr_un cli; /* Unix domain */
socklen_t len;

len = sizeof(cli); /* len is a value */
getpeername(unixfd, (SA *) &cli, &len);
/* len may have changed */
```

The reason that the size changes from an integer to be a po

# Byte Ordering Functions:

converting between the host byte order and the network byte order

	address A+1	address A	
little-endian byte order:	high-order byte	low-order byte	

big-endian byte order:	high-order byte	low-order byte
------------------------	-----------------	----------------

Some machines use the little-endian host byte order while the others use the big-endian. The Internet protocols use the *big-endian* network byte order. Hence, conversion functions should be added in all cases.

```
#include <netinet/in.h>
uint16_t htons(uint16_t host16bitvalue); returns: value in network byte order
uint32_t htonl(uint32_t host32bitvalue); returns: value in network byte order
uint16_t ntohs(uint16_t net16bitvalue); returns: value in host byte order
uint32_t ntohl(uint32_t net32bitvalue); returns: value in host byte order
```



# Byte Manipulation Functions: operating on multibyte fields

From 4.2BSD:

```
#include <strings.h>
void bzero (void *dest, size_t nbytes);
void bcopy (const void *src, void *dest, size_t nbytes);
int bcmp (const void *ptr1, const void *ptr2, size_t nbytes);
                                returns: 0 if equal, nonzero if unequal
```

From ANSI C:

```
#include <string.h>
void *memset (void *dest, int c, size_t len);
void *memcpy (void *dest, const void *src, size_t nbytes);
int memcmp (const void *ptr1, const void *ptr2, size_t nbytes);
                                returns: 0 if equal, nonzero if unequal
```

# Address Conversion Functions

between ASCII strings and network byte ordered binary values

For IPv4 only: ascii and numeric

```
#include <arpa/inet.h>
```

```
int inet_aton (const char *strptr, struct in_addr *addrptr);
```

returns: 1 if string is valid, 0 on error

```
int_addr_t inet_addr (const char *strptr);
```

returns: 32-bit binary IPv4 addr, INADDR\_NONE if error

```
char * inet_ntoa (struct in_addr inaddr);
```

returns: pointer to dotted-decimal string

For IPv4 (AF\_INET) and IPv6 (AF\_INET6): presentation and numeric

```
#include <arpa/inet.h>
```

```
int inet_pton (int family, const char *strptr, void *addrptr);
```

returns: 1 if OK, 0 if invalid presentation, -1 on error

```
const char *inet_ntop (int family, const void *addrptr, char *strptr, size_t len);
```

returns: pointer to result if OK, NULL on error

INET\_ADDRSTRLEN = 16 (for IPv4), INET6\_ADDRSTRLEN = 46 (for IPv6 hex string)

## Example

Even if your system does not yet include support for IPv6, you can start newer functions by replacing calls of the form

```
    foo.sin_addr.s_addr = inet_addr(cp);
```

with

```
    inet_pton(AF_INET, cp, &foo.sin_addr);
```

and replacing calls of the form

```
    ptr = inet_ntoa(foo.sin_addr);
```

with

```
char str[INET_ADDRSTRLEN];
ptr = inet_ntop(AF_INET, &foo.sin_addr, str, sizeof(str));
```

Figure 3.12 shows a simple definition of `inet_pton` that supports only IPv4.  
Figure 3.13 shows a simple version of `inet_ntop` that supports only IPv4.

# Address Conversion Functions (cont.)

Protocol independent functions (not standard system functions):

```
#include "unp.h"
char *sock_ntop (const struct sockaddr *sockaddr, socklen_t addrlen);
    returns: nonnull pointer if OK, NULL on error
```

`sock_bind_wild, sock_cmp_addr, sock_cmp_port, sock_get_port,`  
`sock_ntop_host, sock_set_addr, sock_set_port, sock_set_wild, etc.`

# Stream Socket I/O Functions: preventing callers from having to handle a short count

*read* and *write* return, before reading or writing requested bytes, when the socket buffer limit is reached in the kernel.

When reading from or writing to a stream socket:  
(need to handle EINTR signal)

#include “unp.h”

- `ssize_t readn (int filedes, void *buff, size_t nbytes);`  
                          returns: #bytes read, -1 on error
- `ssize_t writen (int filedes, const void *buff, size_t nbytes);`  
                          returns: #bytes written, -1 on error
- `ssize_t readline (int filedes, void *buff, size_t maxlen);`  
                          returns: #bytes read, -1 on error

Figure 3.15 `readn` function: Read  $n$  bytes from a descriptor.

*lib/readn.c*

```
1 #include      "unp.h"

2 ssize_t          /* Read "n" bytes from a descriptor. */
3 readn(int fd, void *vptr, size_t n)
4 {
5     size_t nleft;
6     ssize_t nread;
7     char   *ptr;

8     ptr = vptr;
9     nleft = n;
10    while (nleft > 0) { 
11        if ((nread = read(fd, ptr, nleft)) < 0) {
12            if (errno == EINTR)
13                nread = 0;           /* and call read() again */
14            else
15                return (-1);
16        } else if (nread == 0) {
17            break;                 /* EOF */
18            nleft -= nread;
19            ptr += nread;
20        }
21    return (n - nleft);           /* return >= 0 */
22 }
```

Figure 3.16 `writen` function: Write  $n$  bytes to a descriptor.

*lib/writen.c*

```
1 #include    "unp.h"

2 ssize_t                      /* Write "n" bytes to a descriptor. */
3 writen(int fd, const void *vptr, size_t n)
4 {
5     size_t nleft;
6     ssize_t nwritten;
7     const char *ptr;

8     ptr = vptr;
9     nleft = n;
10    while (nleft > 0) {
11        if ((nwritten = write(fd, ptr, nleft)) <= 0) {
12            if (nwritten < 0 && errno == EINTR)
13                nwritten = 0; /* and call write() again */
14            else
15                return (-1); /* error */
16        }
17        nleft -= nwritten;
18        ptr += nwritten;
19    }
20    return (n);
21 }
```

Figure 3.17 `readline` function: Read a text line from a descriptor, one byte at a time.

`test/readline1.c`

```
1 #include      "unp.h"  
2 /* PAINFULLY SLOW VERSION -- example only */  
3 ssize_t  
4 readline(int fd, void *vptr, size_t maxlen)  
5 {  
6     ssize_t n, rc;  
7     char    c, *ptr;  
8  
9     ptr = vptr;  
10    for (n = 1; n < maxlen; n++) {  
11        again:  
12            if ((rc = read(fd, &c, 1)) == 1) {  
13                *ptr++ = c;  
14                if (c == '\n') break;           /* newline is stored, like fgets() */  
15            } else if (rc == 0) {  
16                *ptr = 0;  
17                return (n - 1);             /* EOF, n - 1 bytes were read */  
18            } else {  
19                if (errno == EINTR)  
20                    goto again;  
21                return (-1);              /* error, errno set by read() */  
22            }  
23    }  
24  
25    *ptr = 0;                      /* null terminate like fgets() */  
26    return (n);
```

Inefficient but safe

## Figure 3.18 Better version of `readline` function.

`lib/readline.c`

```
1 #include    "unp.h"

2 static int read_cnt;
3 static char *read_ptr;
4 static char read_buf[MAXLINE];

5 static ssize_t
6 my_read(int fd, char *ptr)
7 {

8     if (read_cnt <= 0) {
9         again:
10        if ( (read_cnt = read(fd, read_buf, sizeof(read_buf))) < 0) {
11            if (errno == EINTR)
12                goto again;
13            return (-1);
14        } else if (read_cnt == 0)
15            return (0);
16        read_ptr = read_buf;
17    }

18    read_cnt--;
19    *ptr = *read_ptr++;
20    return (1);
21 }
22 ssize_t
23 readline(int fd, void *unptr, size_t maxlen)
```

```
21
22 ssize_t
23 readline(int fd, void *vptr, size_t maxlen)
24 {
25     ssize_t n, rc;
26     char    c, *ptr;
27
28     ptr = vptr;
29     for (n = 1; n < maxlen; n++) {
30         if ((rc = my_read(fd, &c)) == 1) {  
            *ptr++ = c;  
            if (c == '\n')  
                break; /* newline is stored, like fgets() */  
        } else if (rc == 0) {  
            *ptr = 0;  
            return (n - 1); /* EOF, n - 1 bytes were read */  
        } else  
            return (-1); /* error, errno set by read() */  
    }
38
39     *ptr = 0; /* null terminate like fgets() */
40     return (n);
41 }
42
43 ssize_t
44 readlinebuf(void **vptrptr)
45 {
46     if (read_cnt)
47         *vptrptr = read_ptr;
48     return (read_cnt);
49 }
```

## Efficient but problematic with select()

