

I/O Multiplexing: *select* and *poll*

- Introduction
- I/O models
- *select* function
- Rewrite *str_cli* function
- Supporting batch input with *shutdown* function
- Rewrite concurrent TCP echo server with *select*
- *pselect* function: avoiding signal loss in race condition
- *poll* function: polling more specific conditions than *select*
- Rewrite concurrent TCP echo server with *poll*



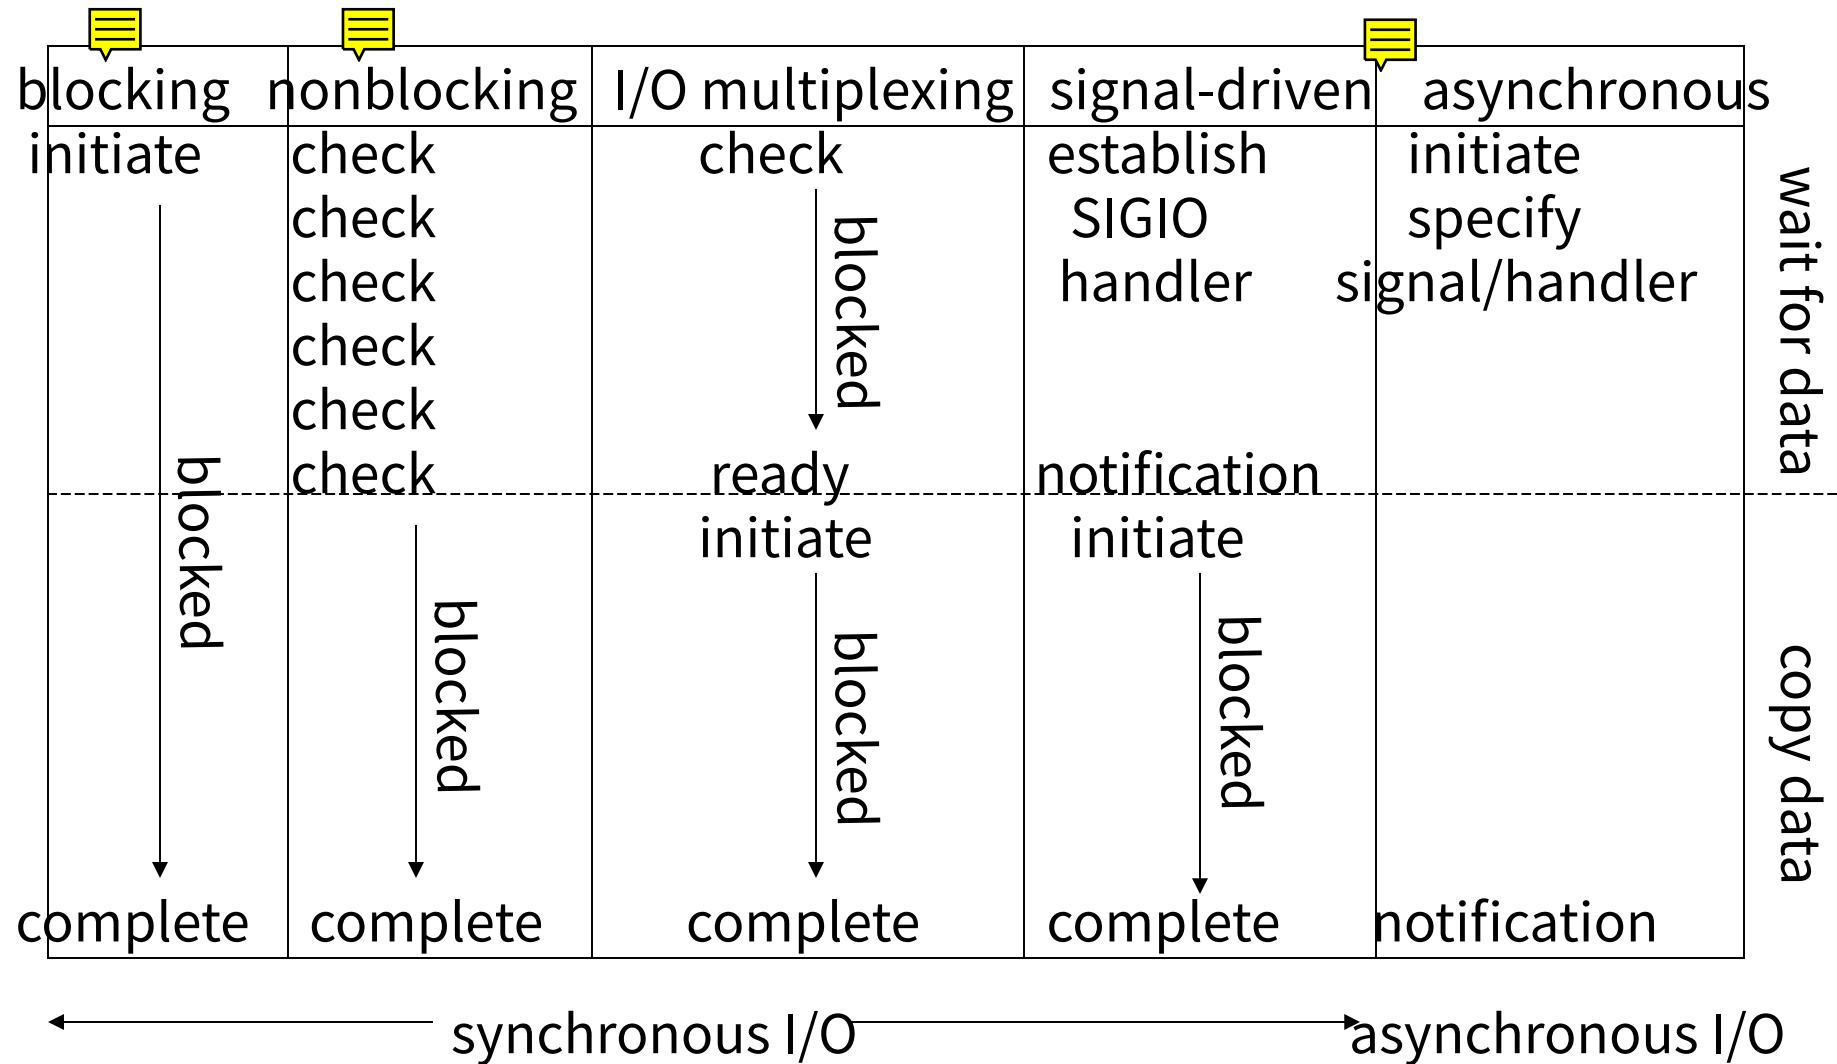
Introduction

- I/O multiplexing: to be notified, by kernel, if one or more I/O conditions are ready.
- Scenarios in networking applications:
 - a client handling multiple descriptors (stdio/socket)
 - a client handling multiple sockets
 - a TCP server handling a listening socket and its connected sockets A small yellow speech bubble icon with three horizontal lines inside, positioned next to the word "connected" in the list.
 - a server handling both TCP and UDP
 - a server handling multiple services and protocols

I/O Models

- Two distinct phases for an input operation:
 - **wait** for data; **copy** data from kernel to user
- Five I/O models:
 - blocking I/O: blocked all the way
 - nonblocking I/O: if no data, immediate returns EWOULDBLOCK
 - I/O multiplexing (*select* and *poll*): blocked separately in wait and copy
 - signal driven I/O (SIGIO): nonblocked in wait but blocked in copy (signaled when I/O can be initiated)
 - asynchronous I/O (*aio_*): nonblocked all the way (signaled when I/O is complete)

Comparison of Five I/O Models



select Function

```
#include <sys/select.h>
#include <sys/time.h>            
int select (int maxfdp1, fd_set *readset, fd_set *writeset, fd_set *exceptset,
            const struct timeval *timeout);
returns: positive count of ready descriptors, 0 on timeout, -1 on error
```

```
struct timeval {          (null: wait forever; 0: do not wait)
    long tv_sec; /*second */
    long tv_usec; /* microsecond */};
```

fd_set --- implementation dependent

four macros:

- void FD_ZERO(fd_set **fdset*); 
- void FD_SET(int *fd*, fd_set **fdset*); 
- void FD_CLR(int *fd*, fd_set **fdset*); 
- int FD_ISSET(int *fd*, fd_set **fdset*); 

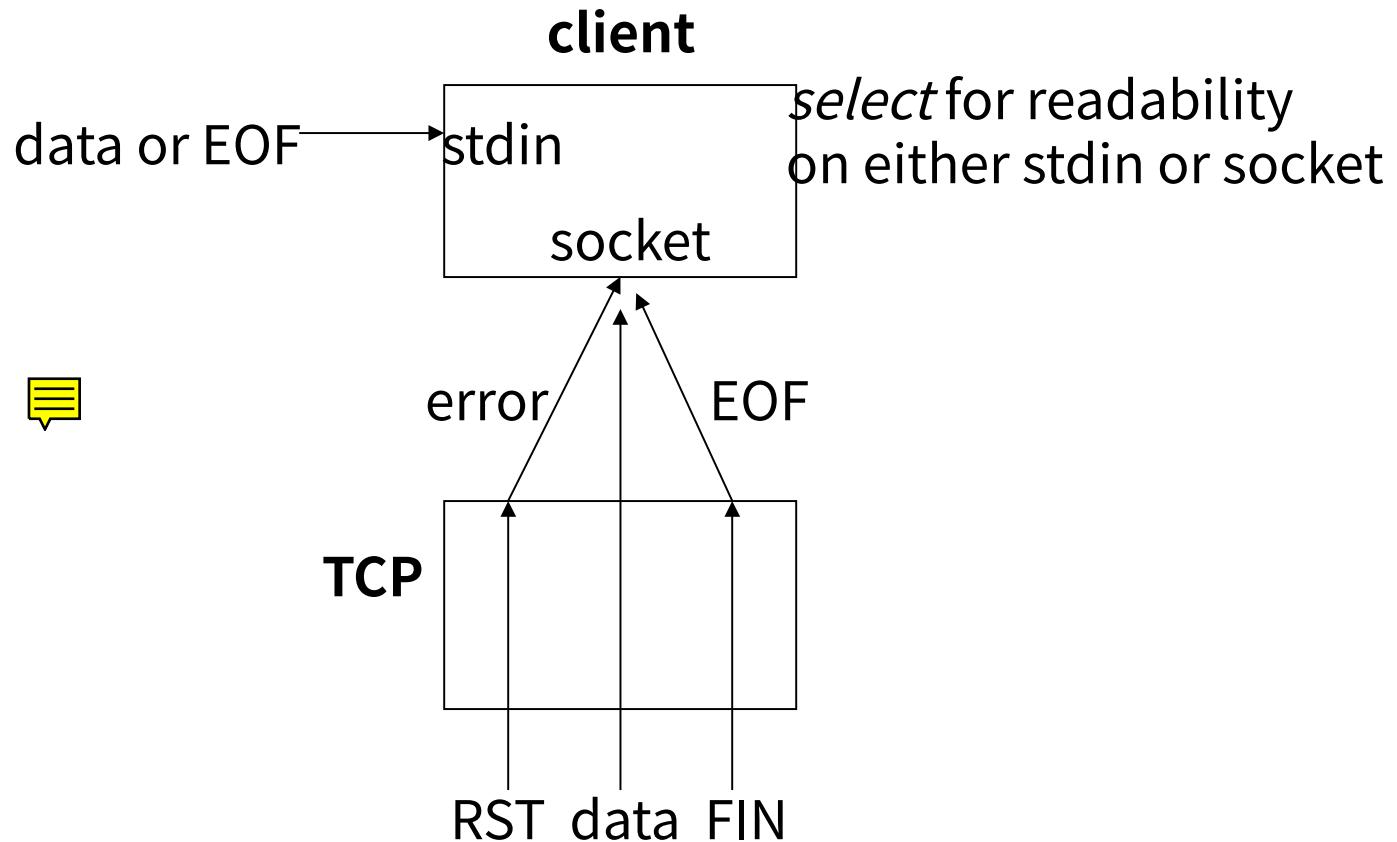
Socket Ready Conditions for *select*

Condition	readable?	writeable?	Exception?
enough data to read read-half closed	x		
new connection ready	x		
writing space available write-half closed		x	x
pending error	x	x	
TCP out-of-band data			x

Low-water mark (enough data/space to read/write in socket receive/send buffer): default is 1/2048, may be set by SO_RCVLOWAT/SO SNDLOWAT socket option

Maximum number of descriptors for select?
Redefine FD_SETSIZE and recompile kernel

Rewrite *str_cli* Function with *select*



Rewrite *str_cli* Function with *select*

```
#include    "unp.h"                                select/strcliselect01.c

void
str_cli(FILE *fp, int sockfd)
{
    int          maxfdp1;
    fd_set      rset;
    char        sendline[MAXLINE], recvline[MAXLINE];

    FD_ZERO(&rset);
    for (;;) {
        FD_SET(fileno(fp), &rset);
        FD_SET(sockfd, &rset);
        maxfdp1 = max(fileno(fp), sockfd) + 1;
        Select(maxfdp1, &rset, NULL, NULL, NULL);
```

select/strcliselect01.c

```
if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
    if (Readline(sockfd, recvline, MAXLINE) == 0)
        err_quit("str_cli: server terminated prematurely
");
    Fputs(recvline, stdout);
}

if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
    if (Fgets(sendline, MAXLINE, fp) == NULL)
        return; /* all done */
    Writen(sockfd, sendline, strlen(sendline));
}
}
```



Supporting Batch Input with *shutdown*

- Stop-and-wait mode (interactive) vs batch mode (redirected stdin/stdout)
- In batch mode, *str_cli* returns right after EOF on input and *main* returns immediately, while leaving server replies unprocessed.
- Solution: In *str_cli*, close write-half of TCP connection, by *shutdown*, while leaving read-half open.

```
#include <sys/socket.h>
int shutdown (int sockfd, int howto);  returns: 0 if OK, -1 on error
howto: SHUT_RD, SHUT_WR, SHUT_RDWR
initiate TCP normal termination regardless of descriptor's reference count
```



Rewrite *str_cli* with *select* and *shutdown*

```
#include    "unp.h"                                select/strcliselect02.c

void
str_cli(FILE *fp, int sockfd)
{
    int      maxfdp1, stdineof;
    fd_set   rset;
    char     sendline[MAXLINE], recvline[MAXLINE];

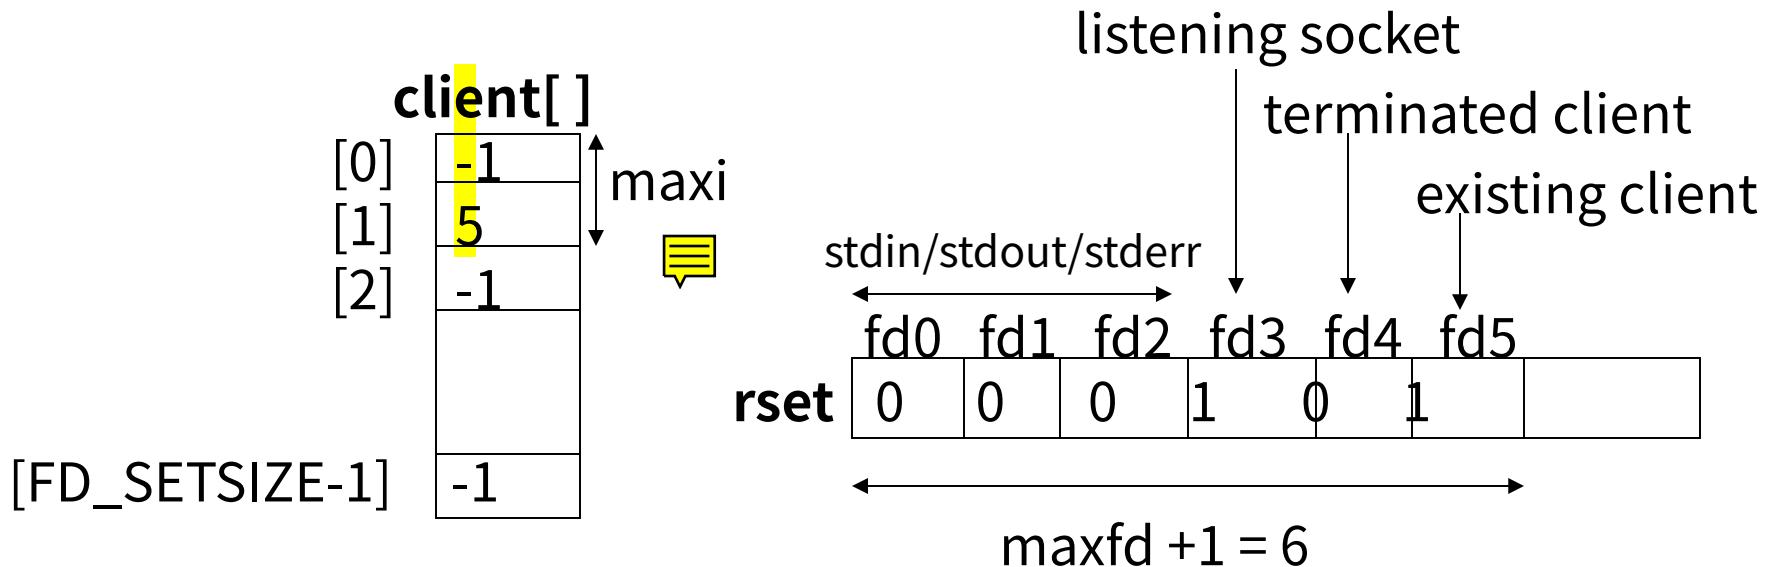
    stdineof = 0;
    FD_ZERO(&rset);
    for (;;) {
        if (stdineof == 0)
            FD_SET(fileno(fp), &rset);
        FD_SET(sockfd, &rset);
        maxfdp1 = max(fileno(fp), sockfd) + 1;
        Select(maxfdp1, &rset, NULL, NULL, NULL);
```

select/strcliselect02.c

```
if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
    if (Readline(sockfd, recvline, MAXLINE) == 0) {
        if (stdineof == 1)
            return;      /* normal termination */
        else
            err_quit("str_cli: server terminated prematurely");
    }
    Fputs(recvline, stdout);
}
if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
    if (Fgets(sendline, MAXLINE, fp) == NULL) {
        stdineof = 1;
        Shutdown(sockfd, SHUT_WR); /* send FIN */
        FD_CLR(fileno(fp), &rset);
        continue;
    }
    Writen(sockfd, sendline, strlen(sendline));
}
}
```

Rewrite TCP Echo Server with *select*

- A single server process using *select* to handle any number of clients
- Need to keep track of the clients by *client[]* (client descriptor array) and *rset* (read descriptor set)



Rewrite TCP Echo Server with *select*

Initialization

tcpcliserv/tcpservselect01.c

```
#include    "unp.h"
int main(int argc, char **argv)
{
    int                  i, maxi, maxfd, listenfd, connfd, sockfd;
    int                  nready, client[FD_SETSIZE];
    ssize_t               n;
    fd_set                rset, allset;
    char                 line[MAXLINE];
    socklen_t              clilen;
    struct sockaddr_in    cliaddr, servaddr;

    listenfd = Socket(AF_INET, SOCK_STREAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family    = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port      = htons(SERV_PORT);
```

Initialization (cont.)

```
Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));  
  
Listen(listenfd, LISTENQ);  
  
maxfd = listenfd; /* initialize */  
maxi = -1; /* index into client[] array */  
for (i = 0; i < FD_SETSIZE; i++)  
    client[i] = -1; /* -1 indicates available entry */  
FD_ZERO(&allset);  
FD_SET(listenfd, &allset);
```

Loop

tcpcliserv/tcpservselect01.c

```
for(;;) {
    rset = allset;      /* structure assignment */
    nready = Select(maxfd+1, &rset, NULL, NULL, NULL);
    if (FD_ISSET(listenfd, &rset)) { /* new client connection */
        clilen = sizeof(cliaddr);
        connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);
        #ifdef NOTDEF
        printf("new client: %s, port %d\n",
               Inet_ntop(AF_INET, &cliaddr.sin_addr, 4, NULL),
               ntohs(cliaddr.sin_port));
        #endif
        for (i = 0; i < FD_SETSIZE; i++)
            if (client[i] < 0) {
                client[i] = connfd; /* save descriptor */
                break;
            }
        if (i == FD_SETSIZE)
            err_quit("too many clients");
        FD_SET(connfd, &allset); /* add new descriptor to set */
        if (connfd > maxfd)
            maxfd = connfd;      /* for select */
        if (i > maxi)
            maxi = i;           /* max index in client[] array */
        if (--nready <= 0)
            continue;           /* no more readable descriptors */
    }
}
```

Loop (cont.)

tcpcliserv/tcpservselect01.c

```
for (i = 0; i <= maxi; i++) { /* check all clients for data */
    if ( (sockfd = client[i]) < 0)
        continue;
    if (FD_ISSET(sockfd, &rset)) {
        if ( (n = Readline(sockfd, line, MAXLINE)) == 0) {
            /* connection closed by client */
            Close(sockfd);
            FD_CLR(sockfd, &allset);
            client[i] = -1;           
        } else
            Writen(sockfd, line, n);
        if (--nready <= 0)
            break;                  /* no more readable descriptors */
    }
}
```

Denial of Service Attacks

- Problem of concurrent TCP echo server with *select* : blocks in a function call, *read* in *readline*, related to a single client
- Attack scenario:
 - a malicious client sends 1 byte of data (other than a newline) and sleep 
 - server hangs until the malicious client either sends a newline or terminates
- Solutions:
 - nonblocking I/O for the listening socket
 - separate thread/process for each client
 - timeout on I/O operations

pselect Function: Avoiding Signal Loss in Race Condition

```
#include <sys/select.h>
#include <signal.h>
#include <time.h>
int pselect (int maxfdp1, fd_set *readset, fd_set *writeset, fd_set *exceptset,
             const struct timespec *timeout, const sigset_t *sigmask);
    returns: count of ready descriptors, 0 on timeout, -1 on error
struct timespec {
                           time_t tv_sec; /* seconds */
                           long tv_nsec; /* nanosecond */};

if (intr_flag)
    handle_intr(); /* handle signal */
if ((nready = select ( ... )) < 0) {
    if (errno == EINTR) {
        if (intr_flag)
            handle_intr();
    }
    ....
}
signal lost if select blocks forever
```

} } }

```
           sigemptyset (&zeromask);
           sigemptyset (&newmask);
           sigaddset (&newmask, SIGINT);
           sigprocmask (SIG_BLOCK, &newmask, &oldmask);
           if (intr_flag)
               handle_intr();
           if ( (nready = pselect ( ... , &zeromask)) < 0) {
               if (errno == EINTR) {
                   if (intr_flag)
                       handle_intr();
               }
           }
       .... }
```

☰ poll Function: polling more specific conditions than select

```
#include <poll.h>
int poll (struct pollfd *fdarray, unsigned long ndfs, int timeout);
    returns: count of ready descriptors, 0 on timeout, -1 on error
struct pollfd {
    int fd;          /* a descriptor to poll */
    short events;   /* events of interest on fd, value argument */
    short revents;  /* events that occurred on fd, result argument */
};
```

Constant	events	revents	Description
POLLIN	x	x	normal or priority band to read
POLLRDNORM	x	x	normal data to read
POLLRDBAND	x	x	priority band data to read
POLLPRI	x	x	high-priority data to read
POLLOUTx	x		normal data to write
POLLWRNORM	x	x	normal data to write
POLLWRBAND	x	x	priority band data to write
POLLERR	x		error occurred
POLLHUP		x	hangup occurred
POLLNVAL		x	descriptor is not an open file

Concurrent TCP Echo Server with *poll*

- When using *select*, the server maintains array *client[]* and descriptor set *rset*. When using *poll*, the server maintains array *client* of *pollfd* structures.
- Program flow:
 - allocate array of *pollfd* structures
 - initialize (listening socket: first entry in *client*)
(set *POLLRDNORM* in *events*)
 - call *poll*; check for new connection
(check, in *revents*, and set, in *events*, *POLLRDNORM*)
 - check for data on an existing connection
(check *POLLRDNORM* or *POLLERR* in *revents*)

Rewrite Concurrent TCP Echo Server with *poll*

Initialization

tcpcliserv/tcpservpoll01.c

```
#include    "unp.h"
#include    <limits.h>      /* for OPEN_MAX */
int main(int argc, char **argv)
{
    int          i, maxi, listenfd, connfd, sockfd;
    int          nready;
    ssize_t       n;
    char         line[MAXLINE];
    socklen_t     clilen;
    struct pollfd   client[OPEN_MAX];
    struct sockaddr_in  cliaddr, servaddr;
    listenfd = Socket(AF_INET, SOCK_STREAM, 0);
    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family   = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port     = htons(SERV_PORT);
```

Initialization (cont.)

```
Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));
```

```
Listen(listenfd, LISTENQ);
```

```
client[0].fd = listenfd;
client[0].events = POLLRDNORM;
for (i = 1; i < OPEN_MAX; i++)
    client[i].fd = -1;           /* -1 indicates available entry */
maxi = 0;                      /* max index into client[] array */
```

Loop

```

for (;;) {
    nready = Poll(client, maxi+1, INFTIM);

    if (client[0].revents & POLLRDNORM) { /* new client connection */
        clilen = sizeof(cliaddr);
        connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);
        #ifdef NOTDEF
        printf("new client: %s\n", Sock_ntop((SA *) &cliaddr, clilen));
        #endif
        for (i = 1; i < OPEN_MAX; i++)
            if (client[i].fd < 0) {
                client[i].fd = connfd; /* save descriptor */
                break;
            }
        if (i == OPEN_MAX)
            err_quit("too many clients");
        client[i].events = POLLRDNORM;
        if (i > maxi)

            maxi = i;                  /* max index in client[] array */
        if (--nready <= 0)
            continue;                 /* no more readable descriptors */
    }
}

```

Loop (cont.)

tcpcliserv/tcpservpoll01.c

```
for (i = 1; i <= maxi; i++) { /* check all clients for data */
    if ( (sockfd = client[i].fd) < 0)
        continue;
    if (client[i].revents & (POLLRDNORM | POLLERR)) {
        if ( (n = readline(sockfd, line, MAXLINE)) < 0) {
            if (errno == ECONNRESET) {
                /* connection reset by client */
                #ifdef NOTDEF
                printf("client[%d] aborted connection\n", i);
                #endif
                Close(sockfd);
                client[i].fd = -1;
            } else
                err_sys("readline error");
        } else if (n == 0) {
            /* connection closed by client */
            #ifdef NOTDEF
            printf("client[%d] closed connection\n", i);
            #endif
            Close(sockfd);
            client[i].fd = -1;
        } else
            Writen(sockfd, line, n);
        if (--nready <= 0)
            break;           /* no more readable descriptors */
    }
}
```