

Nonblocking I/O

- Nonblocking reads and writes
 - Default mode is blocking
 - TCP v.s. UDP
- Nonblocking accept
- Nonblocking connect

Nonblocking Reads and Writes overlapping I/O in str_cli of echo client

- Set descriptors nonblocking using fcntl
- initialize buffer pointers of “to” and “fr” buffers
- main loop to call select
 - specify interested descriptors
 - call select
 - read from standard input
 - handle nonblocking error
 - read returns end-of-file
 - read returns data
 - read from socket
 - write OK
 - write to socket

Buffers Enabling Overlapped Nonblocking I/O in str_cli

Figure 16.1. Buffer containing data from standard input going to the socket.

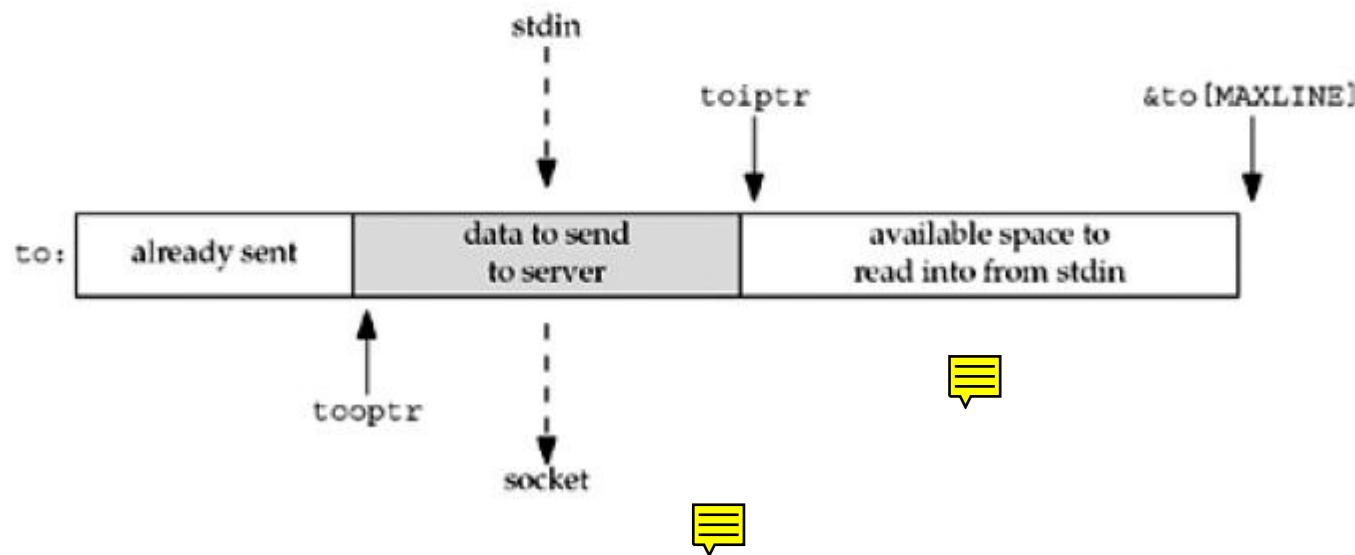


Figure 16.2. Buffer containing data from the socket going to standard output.

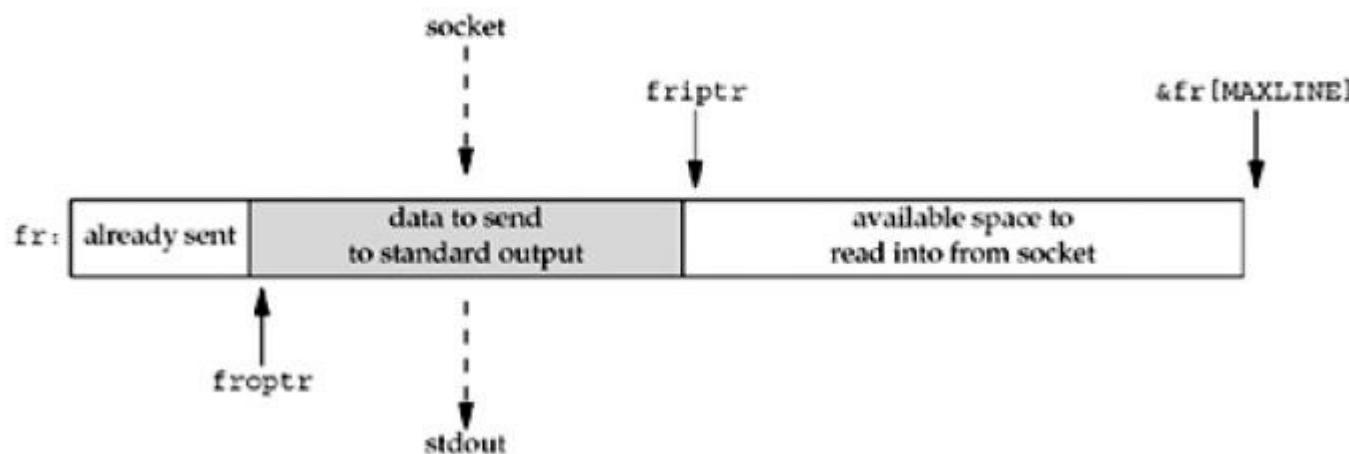


Figure 16.3 `str_cli` function, first part: initializes and calls `select`.

nonblock/strclononb.c

```
1 #include      "unp.h"
2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     int      maxfdp1, val, stdineof;
6     ssize_t n, nwritten;
7     fd_set  rset, wset;
8     char    to[MAXLINE], fr[MAXLINE];
9     char    *toiptr, *tooptr, *friptr, *froptr;
10
11    val = Fcntl(sockfd, F_GETFL, 0);
12    Fcntl(sockfd, F_SETFL, val | O_NONBLOCK); 
13
14    val = Fcntl(STDIN_FILENO, F_GETFL, 0);
15    Fcntl(STDIN_FILENO, F_SETFL, val | O_NONBLOCK);
16
17    val = Fcntl(STDOUT_FILENO, F_GETFL, 0);
18    Fcntl(STDOUT_FILENO, F_SETFL, val | O_NONBLOCK); 
19
20    toiptr = tooptr = to;           /* initialize buffer pointers */
21    friptr = froptr = fr;
22    stdineof = 0; 
23
24    maxfdp1 = max(max(STDIN_FILENO, STDOUT_FILENO), sockfd) + 1;
25    for ( ; ; ) {
26        FD_ZERO(&rset);
```

```
10     val = Fcntl(sockfd, F_GETFL, 0);
11     Fcntl(sockfd, F_SETFL, val | O_NONBLOCK);

12     val = Fcntl(STDIN_FILENO, F_GETFL, 0);
13     Fcntl(STDIN_FILENO, F_SETFL, val | O_NONBLOCK);

14     val = Fcntl(STDOUT_FILENO, F_GETFL, 0);
15     Fcntl(STDOUT_FILENO, F_SETFL, val | O_NONBLOCK);

16     toiptr = tooptr = to;           /* initialize buffer pointers */
17     friptr = froptr = fr;
18     stdineof = 0;

19     maxfdpl = max(max(STDIN_FILENO, STDOUT_FILENO), sockfd) + 1;
20     for ( ; ; ) {
21         FD_ZERO(&rset);
22         FD_ZERO(&wset);
23         if (stdineof == 0 && toiptr < &to[MAXLINE])
24             FD_SET(STDIN_FILENO, &rset);      /* read from stdin */
25         if (friptr < &fr[MAXLINE])
26             FD_SET(sockfd, &rset);    /* read from socket */
27         if (tooptr != toiptr)
28             FD_SET(sockfd, &wset);   /* data to write to socket */
29         if (froptr != friptr)
30             FD_SET(STDOUT_FILENO, &wset); /* data to write to stdout */

31     Select(maxfdpl, &rset, &wset, NULL, NULL);
```

Figure 16.4 str_cli function, second part: reads from standard input or socket.

nonblock/strclintonb.c

```
32     if (FD_ISSET(STDIN_FILENO, &rset)) {  
33         if ((n = read(STDIN_FILENO, toiptr, &to[MAXLINE] - toiptr)) < 0) {  
34             if (errno != EWOULDBLOCK)  
35                 err_sys("read error on stdin");  
  
36         } else if (n == 0) {  
37             fprintf(stderr, "%s: EOF on stdin\n", gf_time());  
38             stdineof = 1; /* all done with stdin */  
39             if (tooptr == toiptr)  
40                 Shutdown(sockfd, SHUT_WR); /* send FIN */  
  
41         } else {  
42             fprintf(stderr, "%s: read %d bytes from stdin\n", gf_time(),  
43                     n);  
44             toiptr += n; /* # just read */  
45             FD_SET(sockfd, &wset); /* try and write to socket below */  
46         }  
47     }  
  
48     if (FD_ISSET(sockfd, &rset)) {  
49         if ((n = read(sockfd, friptr, &fr[MAXLINE] - friptr)) < 0) {  
50             if (errno != EWOULDBLOCK)  
51                 err_sys("read error on socket");  
  
52         } else if (n == 0) {  
53             if (stdineof) /* no more EOF on standard input */  
54                 err_sys("socket closed");  
55         }  
56     }  
57 }
```

```
39         if (tooptr == toiptr)
40             Shutdown(sockfd, SHUT_WR); /* send FIN */
41
42     } else {
43         fprintf(stderr, "%s: read %d bytes from stdin\n", gf_time(),
44                 n);
45         toiptr += n; /* # just read */
46         FD_SET(sockfd, &wset); /* try and write to socket below */
47     }
48
49 if (FD_ISSET(sockfd, &rset)) {
50     if ((n = read(sockfd, friptr, &fr[MAXLINE] - friptr)) < 0) {
51         if (errno != EWOULDBLOCK)
52             err_sys("read error on socket");
53
54     } else if (n == 0) {
55         fprintf(stderr, "%s: EOF on socket\n", gf_time());
56         if (stdineof)
57             return; /* normal termination */
58     } else {
59         fprintf(stderr, "%s: read %d bytes from socket\n",
60                 gf_time(), n);
61         friptr += n; /* # just read */
62         FD_SET(STDOUT_FILENO, &wset); /* try and write below */
63     }
64 }
```

Figure 16.5 `str_cli` function, third part: writes to standard output or socket.

nonblock/strclintonb.c

```
65     if (FD_ISSET(STDOUT_FILENO, &wset) && ((n = friptr - froptr) > 0)) {
66         if ( (nwritten = write(STDOUT_FILENO, froptr, n)) < 0) {
67             if (errno != EWOULDBLOCK)
68                 err_sys("write error to stdout");
69
70         } else {
71             fprintf(stderr, "%s: wrote %d bytes to stdout\n",
72                     gf_time(), nwritten);
73             froptr += nwritten; /* # just written */
74             if (froptr == friptr)
75                 froptr = friptr = fr; /* back to beginning of buffer */
76         }
77
78     if (FD_ISSET(sockfd, &wset) && ((n = toiptr - tooptr) > 0)) {
79         if ( (nwritten = write(sockfd, tooptr, n)) < 0) {
80             if (errno != EWOULDBLOCK)
81                 err_sys("write error to socket");
82
83         } else {
84             fprintf(stderr, "%s: wrote %d bytes to socket\n",
85                     gf_time(), nwritten);
86             tooptr += nwritten; /* # just written */
87             if (toiptr == friptr) {
88                 toiptr = friptr = fr;
89             }
90         }
91     }
92 }
```

```
71             gf_time(), nwritten);
72     froptr += nwritten; /* # just written */
73     if (froptr == friptr)
74         froptr = friptr = fr; /* back to beginning of buffer */
75     }
76 }

77 if (FD_ISSET(sockfd, &wset) && ((n = toiptr - tooptr) > 0)) {
78     if ( (nwritten = write(sockfd, tooptr, n)) < 0) {
79         if (errno != EWOULDBLOCK)
80             err_sys("write error to socket");

81     } else {
82         fprintf(stderr, "%s: wrote %d bytes to socket\n",
83                 gf_time(), nwritten);
84         tooptr += nwritten; /* # just written */
85         if (tooptr == toiptr) {
86             toiptr = tooptr = to; /* back to beginning of buffer */
87             if (stdineof)
88                 Shutdown(sockfd, SHUT_WR); /* send FIN */
89         }
90     }
91 }
92 }
93 }
```

Figure 16.6 `gf_time` function: returns pointer to time string.

lib/gf_time.c

```
1 #include      "unp.h"
2 #include      <time.h>

3 char *
4 gf_time(void)
5 {
6     struct timeval tv;
7     static char str[30];
8     char *ptr;

9     if (gettimeofday(&tv, NULL) < 0)
10        err_sys("gettimeofday error");

11    ptr = ctime(&tv.tv_sec);
12    strcpy(str, &ptr[11]);
13    /* Fri Sep 13 00:00:00 1986\n\0 */
14    /* 0123456789012345678901234 5 */
15    snprintf(str + 8, sizeof(str) - 8, ".%06ld", tv.tv_usec);

16    return (str);
17 }
```

An Usage Example

```
solaris % tcpdump -w tcpd tcp and port 7
```

We then run our TCP client on this host, specifying the server on the host linux.

```
solaris % tcpcli02 192.168.1.10 < 2000.lines > out 2> diag
```

Standard input is the file 2000.lines, the same file we used with Figure 6.13. Standard output is sent to the file out, and standard error is sent to the file diag. On completion, we run

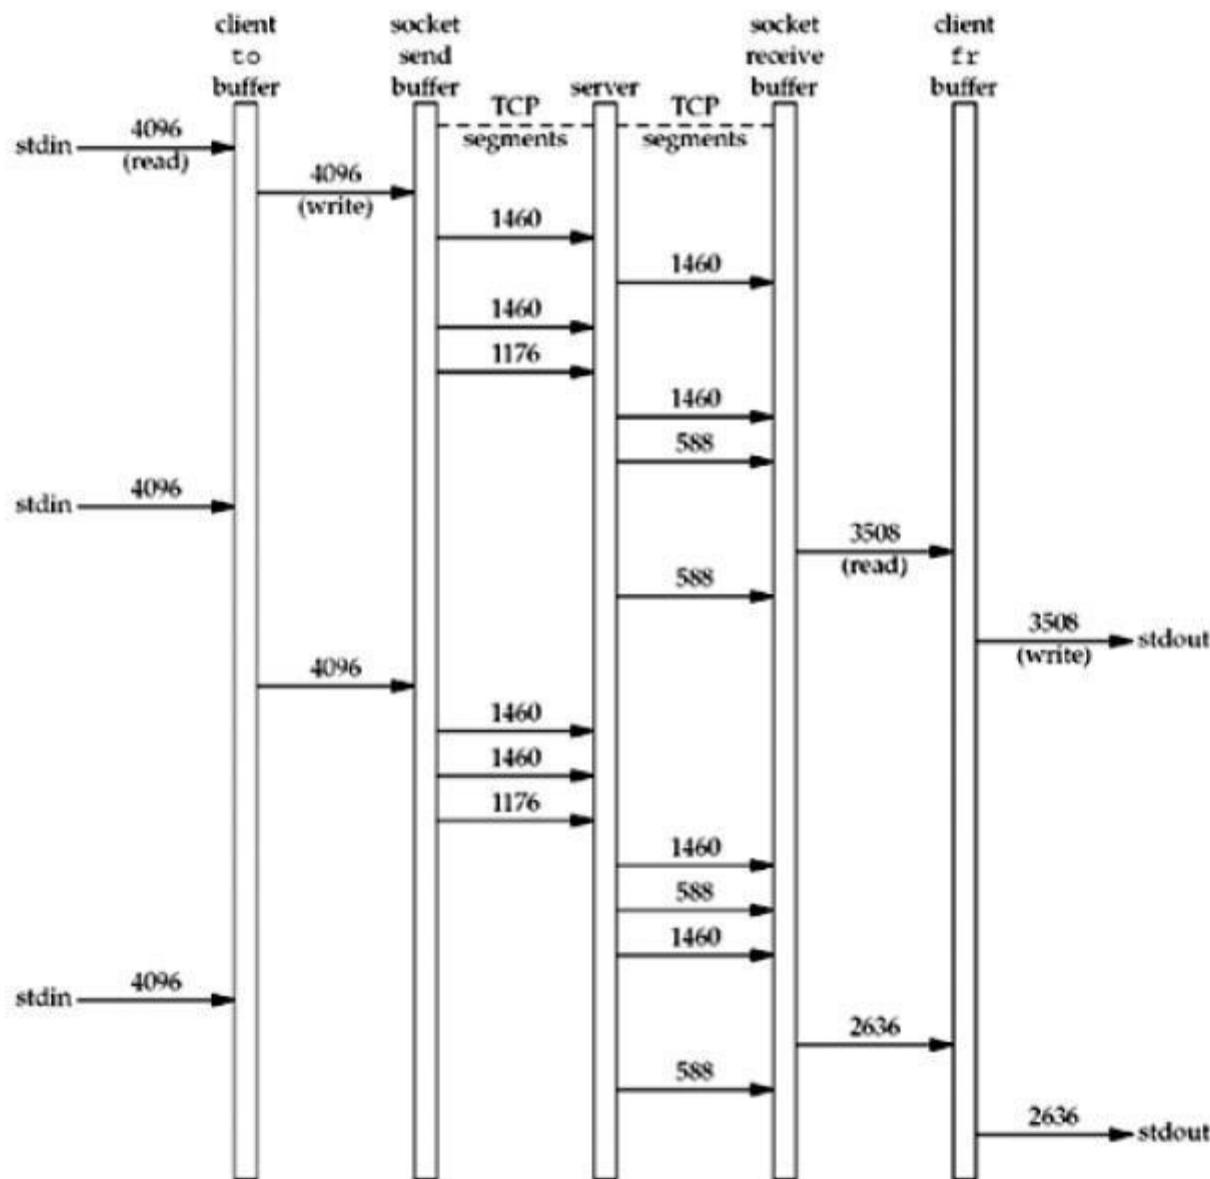
```
solaris % diff 2000.lines out
```

Figure 16.7 Sorted output from tcpdump and diagnostic output.

```
solaris % tcpdump -r tcpd -N | sort diag -
10:18:34.486392 solaris.33621 > linux.echo: S 1802738644:1802738644(0)
                                         win 8760 <mss 1460>
10:18:34.488278 linux.echo > solaris.33621: S 3212986316:3212986316(0)
                                         ack 1802738645 win 8760 <mss 1460>
10:18:34.488490 solaris.33621 > linux.echo: . ack 1 win 8760

10:18:34.491482: read 4096 bytes from stdin
10:18:34.518663 solaris.33621 > linux.echo: P 1:1461(1460) ack 1 win 8760
10:18:34.519016: wrote 4096 bytes to socket
10:18:34.528529 linux.echo > solaris.33621: P 1:1461(1460) ack 1461 win 8760
10:18:34.528785 solaris.33621 > linux.echo: . 1461:2921(1460) ack 1461 win 8760
10:18:34.528900 solaris.33621 > linux.echo: P 2921:4097(1176) ack 1461 win 8760
10:18:34.528958 solaris.33621 > linux.echo: . ack 1461 win 8760
10:18:34.536193 linux.echo > solaris.33621: . 1461:2921(1460) ack 4097 win 8760
10:18:34.536697 linux.echo > solaris.33621: P 2921:3509(588) ack 4097 win 8760
10:18:34.544636: read 4096 bytes from stdin
10:18:34.568505: read 3508 bytes from socket
10:18:34.580373 solaris.33621 > linux.echo: . ack 3509 win 8760
10:18:34.582244 linux.echo > solaris.33621: P 3509:4097(588) ack 4097 win 8760
10:18:34.593354: wrote 3508 bytes to stdout
10:18:34.617272 solaris.33621 > linux.echo: P 4097:5557(1460) ack 4097 win 8760
10:18:34.617610 solaris.33621 > linux.echo: P 5557:7017(1460) ack 4097 win 8760
10:18:34.617908 solaris.33621 > linux.echo: P 7017:8193(1176) ack 4097 win 8760
10:18:34.618062: wrote 4096 bytes to socket
10:18:34.623310 linux.echo > solaris.33621: . ack 8193 win 8760
10:18:34.626129 linux.echo > solaris.33621: . 4097:5557(1460) ack 8193 win 8760
10:18:34.626339 solaris.33621 > linux.echo: . ack 5557 win 8760
10:18:34.626611 linux.echo > solaris.33621: P 5557:6145(588) ack 8193 win 8760
10:18:34.628396 linux.echo > solaris.33621: . 6145:7605(1460) ack 8193 win 8760
10:18:34.643524: read 4096 bytes from stdin
10:18:34.667305: read 2636 bytes from socket
10:18:34.670324 solaris.33621 > linux.echo: . ack 7605 win 8760
10:18:34.672221 linux.echo > solaris.33621: P 7605:8193(588) ack 8193 win 8760
10:18:34.691039: wrote 2636 bytes to stdout
```

Figure 16.8. Timeline of nonblocking example.



Is the effort worth it at all?

- The nonblocking version is nontrivial: about 135 lines of code
- The version using select with blocking I/O in Fig. 6.13 has only 36 lines of code.
- The original stop-and-wait version in Fig. 5.5 has only 12 lines of code.

Figure 6.13 `str_cli` function using `select` that handles EOF correctly.

`select/strcli/select02.c`

```
1 #include    "unp.h"
2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     int      maxfdp1, stdineof;
6     fd_set  rset;
7     char    buf[MAXLINE];
8     int      n;

9     stdineof = 0;
10    FD_ZERO(&rset);
11    for ( ; ; ) {
12        if (stdineof == 0)
13            FD_SET(fileno(fp), &rset);
14        FD_SET(sockfd, &rset);
15        maxfdp1 = max(fileno(fp), sockfd) + 1;
16        Select(maxfdp1, &rset, NULL, NULL, NULL);

17        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
18            if ( (n = Read(sockfd, buf, MAXLINE)) == 0) {
19                if (stdineof == 1)
20                    return; /* normal termination */
21                else
22                    err_quit("str_cli: server terminated prematurely");
23            }
24            Write(fileno(stdout), buf, n);
25        }
26        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
27            if ( (n = Read(fileno(fp), buf, MAXLINE)) == 0) {
28                stdineof = 1;
29                Shutdown(sockfd, SHUT_WR); /* send FIN */
30                FD_CLR(fileno(fp), &rset);
31                continue;
32            }
33            Writen(sockfd, buf, n);
34        }
35    }
36 }
```

Select() with blocking I/O (40 lines)

Figure 5.5 `str_cli` function: client processing loop.

lib/str_cli.c

The original stop-and-wait version (12 lines)

```
1 #include    "unp.h"
2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     char      sendline[MAXLINE], recvline[MAXLINE];
6
7     while (Fgets(sendline, MAXLINE, fp) != NULL) {
8
9         Writen(sockfd, sendline, strlen (sendline));
10
11        if (Readline(sockfd, recvline, MAXLINE) == 0)
12            err_quit("str_cli: server terminated prematurely");
13
14        Fputs(recvline, stdout);
15    }
16 }
```



A Simpler Version Using Two Processes

Figure 16.9. `str_cli` function using two processes.

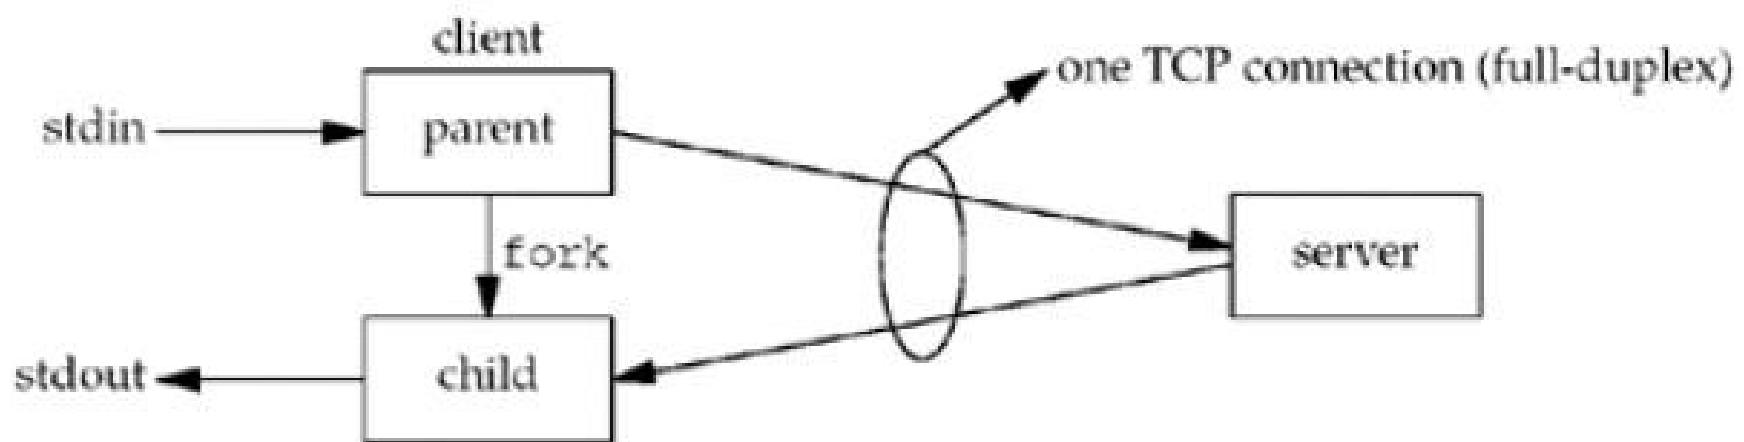


Figure 16.10 Version of str_cli function that uses fork.

nonblock/strclifork.c

```
1 #include      "unp.h"

2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     pid_t      pid;
6     char      sendline[MAXLINE], recvline[MAXLINE];

7     if ((pid = Fork()) == 0) { /* child: server -> stdout */
8         while (Readline(sockfd, recvline, MAXLINE) > 0)
9             Fputs(recvline, stdout);

10        kill(getppid(), SIGTERM); /* in case parent still running */
11        exit(0);
12    }

13    /* parent: stdin -> server */
14    while (Fgets(sendline, MAXLINE, fp) != NULL)
15        Writen(sockfd, sendline, strlen(sendline));

16    Shutdown(sockfd, SHUT_WR); /* EOF on stdin, send FIN */
17    pause();
18    return;
19 }
```

Performance Comparison

- 354.0 sec, stop-and-wait (Figure 5.5)
- 12.3 sec, select and blocking I/O (Figure 6.13)
- 6.9 sec, nonblocking I/O (Figure 16.3)
- 8.7 sec, fork (Figure 16.10)
- 8.5 sec, threaded version (Figure 26.2)

Your choice?

Nonblocking connect: Daytime Client

- Set socket nonblocking by *fcntl*
- Overlap processing with connection establishment
- Check for immediate completion
- Call *select*
- Handle timeouts
- Check for readability or writability
- Turn off nonblocking and return

Figure 16.11 Issue a nonblocking connect.

lib/connect_nonb.c

```
1 #include      "unp.h"

2 int
3 connect_nonb(int sockfd, const SA *saptr, socklen_t salen, int nsec)
4 {
5     int      flags, n, error;
6     socklen_t len;
7     fd_set rset, wset;
8     struct timeval tval;
9     flags = Fcntl(sockfd, F_GETFL, 0);
10    Fcntl(sockfd, F_SETFL, flags | O_NONBLOCK);
11    error = 0;
12    if ( (n = connect(sockfd, saptr, salen)) < 0)
13        if (errno != EINPROGRESS)
14            return (-1);
15    /* Do whatever we want while the connect is taking place. */

16    if (n == 0)
17        goto done;                                /* connect completed immediately */

18    FD_ZERO(&rset);
19    FD_SET(sockfd, &rset);
20    wset = rset;
21    tval.tv_sec = nsec;
22    tval.tv_usec = 0;

23    if ( (n = Select(sockfd + 1, &rset, &wset, NULL,
24                      nsec ? &tval : NULL)) == 0) {
25        close(sockfd);                         /* timeout */
26        errno = ETIMEDOUT;
27        return (-1);
28    }
```

```
21     tval.tv_sec = nsec,
22     tval.tv_usec = 0;
23
24     if ( (n = Select(sockfd + 1, &rset, &wset, NULL,
25                         nsec ? &tval : NULL)) == 0) {
26         close(sockfd);           /* timeout */
27         errno = ETIMEDOUT;
28         return (-1);
29     }
30     if (FD_ISSET(sockfd, &rset) || FD_ISSET(sockfd, &wset)) {
31         len = sizeof(error);
32         if (getsockopt(sockfd, SOL_SOCKET, SO_ERROR, &error, &len) < 0)
33             return (-1);        /* Solaris pending error */
34     } else
35         err_quit("select error: sockfd not set");
36
37     done:
38     Fcntl(sockfd, F_SETFL, flags); /* restore file status flags */
39
40     if (error) {
41         close(sockfd);          /* just in case */
42         errno = error;
43         return (-1);
44     }
45     return (0);
46 }
```

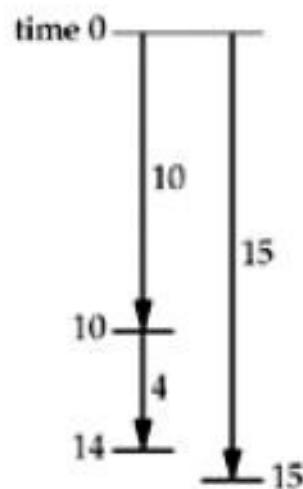
Nonblocking connect: Web Client

- Do you know why Netscape web browser beat Mosaic web browser in 1995?
- It uses the parallel download technique.

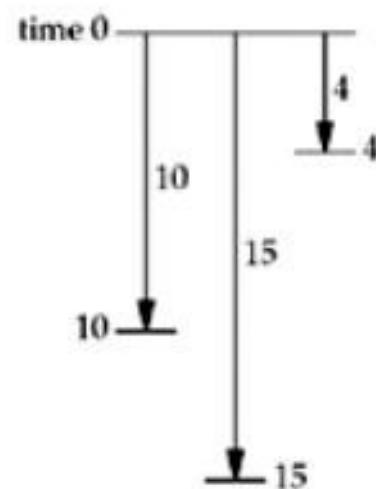
Figure 16.12. Establishing multiple connections in parallel.



three connections
done serially



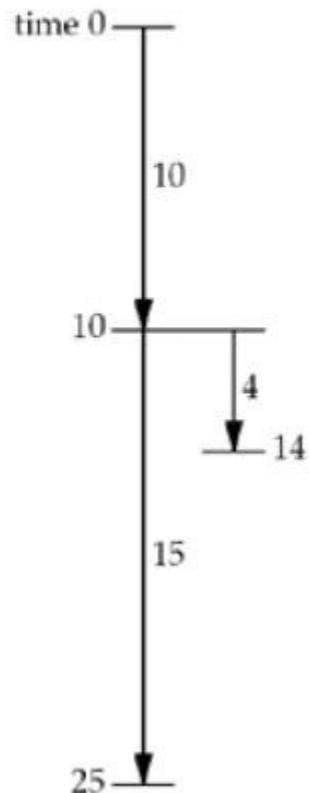
three connections
done in parallel;
maximum of two
connections at a time



three connections
done in parallel;
maximum of three
connections at a time

The first web page download can only be done serially

Figure 16.13. Complete first connection, then multiple connections in parallel.



An Usage Example of a Parallel-Download Web Browser

```
solaris % web 3 www.foobar.com / image1.gif image2.gif \
image3.gif image4.gif image5.gif \
image6.gif image7.gif
```

Figure 16.14 `web.h` header.

nonblock/web.h

```
1 #include      "unp.h"

2 #define MAXFILES    20
3 #define SERV        "80"           /* port number or service name */

4 struct file {
5     char   *f_name;             /* filename */
6     char   *f_host;            /* hostname or IPv4/IPv6 address */
7     int    f_fd;               /* descriptor */
8     int    f_flags;            /* F_xxx below */
9 } file[MAXFILES];

10 #define F_CONNECTING   1           /* connect() in progress */
11 #define F_READING      2           /* connect() complete; now reading */
12 #define F_DONE          4           /* all done */

13 #define GET_CMD         "GET %s HTTP/1.0\r\n\r\n"

14             /* globals */
15 int      nconn,  nfiles, nlefttoconn, nlefttoread, maxfd;
16 fd_set  rset,  wset;

17             /* function prototypes */
18 void    home_page(const char *, const char *);
19 void    start_connect(struct file *);
20 void    write_get_cmd(struct file );
```

Figure 16.15 First part of simultaneous connect: globals and start of main.

nonblock/web.c

```
1 #include      "web.h"

2 int
3 main(int argc, char **argv)
4 {
5     int      i, fd, n, maxnconn, flags, error;
6     char    buf[MAXLINE];
7     fd_set  rs, ws;

8     if (argc < 5)
9         err_quit("usage: web <#conns> <hostname> <homepage> <file1> ...");
10    maxnconn = atoi(argv[1]);

11    nfiles = min(argc - 4, MAXFILES);
12    for (i = 0; i < nfiles; i++) {
13        file[i].f_name = argv[i + 4];
14        file[i].f_host = argv[2];
15        file[i].f_flags = 0;
16    }
17    printf("nfiles = %d\n", nfiles);

18    home_page(argv[2], argv[3]);

19    FD_ZERO(&rset);
20    FD_ZERO(&wset);
21    maxfd = -1;
22    nlefttoread = nlefttoconn = nfiles;
23    nconn = 0;
```

Figure 16.16 *home_page* **function.**

nonblock/home_page.c

```
1 #include      "web.h"

2 void
3 home_page(const char *host, const char *fname)
4 {
5     int      fd, n;
6     char    line[MAXLINE];
7     fd = Tcp_connect(host, SERV); /* blocking connect() */

8     n = sprintf(line, sizeof(line), GET_CMD, fname);
9     Write(fd, line, n);

10    for ( ; ; ) {
11        if ( (n = Read(fd, line, MAXLINE)) == 0)
12            break;                  /* server closed connection */

13        printf("read %d bytes of home page\n", n);
14        /* do whatever with data */
15    }
16    printf("end-of-file on home page\n");
17    Close(fd);
18 }
```

Figure 16.17 Initiate nonblocking connect.

nonblock/start_connect.c

```
1 #include      "web.h"

2 void
3 start_connect(struct file *fptr)
4 {
5     int      fd, flags, n;
6     struct addrinfo *ai;

7     ai = Host_serv(fptr->f_host, SERV, 0, SOCK_STREAM);

8     fd = Socket(ai->ai_family, ai->ai_socktype, ai->ai_protocol);
9     fptr->f_fd = fd;
10    printf("start_connect for %s, fd %d\n", fptr->f_name, fd);

11    /* Set socket nonblocking */
12    flags = Fcntl(fd, F_GETFL, 0);
13    Fcntl(fd, F_SETFL, flags | O_NONBLOCK);

14    /* Initiate nonblocking connect to the server. */
15    if ( (n = connect(fd, ai->ai_addr, ai->ai_addrlen)) < 0) {
16        if (errno != EINPROGRESS)
17            err_svs("nonblocking connect error");
```

```
11     /* Set socket nonblocking */
12 flags = Fcntl(fd, F_GETFL, 0);
13 Fcntl(fd, F_SETFL, flags | O_NONBLOCK);

14     /* Initiate nonblocking connect to the server. */
15 if ( (n = connect(fd, ai->ai_addr, ai->ai_addrlen)) < 0) {
16     if (errno != EINPROGRESS)
17         err_sys("nonblocking connect error");
18     fptr->f_flags = F_CONNECTING;
19     FD_SET(fd, &rset);      /* select for reading and writing */
20     FD_SET(fd, &wset);
21     if (fd > maxfd)
22         maxfd = fd;

23 } else if (n >= 0)           /* connect is already done */
24     write_get_cmd(fptr);      /* write() the GET command */
25 }
```

Figure 16.18 Send an HTTP GET command to the server.

nonblock/write_get_cmd.c

```
1 #include      "web.h"

2 void
3 write_get_cmd(struct file *fptr)
4 {
5     int      n;
6     char    line[MAXLINE];

7     n = snprintf(line, sizeof(line), GET_CMD, fptr->f_name);
8     written(fptr->f_fd, line, n);
9     printf("wrote %d bytes for %s\n", n, fptr->f_name);

10    fptr->f_flags = F_READING; /* clears F_CONNECTING */
11    FD_SET(fptr->f_fd, &rset); /* will read server's reply */
12    if (fptr->f_fd > maxfd)
13        maxfd = fptr->f_fd;
14 }
```

Figure 16.19 Main loop of `main` function.

nonblock/web.c

```
24     while (nlefttoread > 0) {
25         while (nconn < maxnconn && nlefttoconn > 0) {
26             /* find a file to read */
27             for (i = 0; i < nfiles; i++)
28                 if (file[i].f_flags == 0)
29                     break;
30             if (i == nfiles)
31                 err_quit("nlefttoconn = %d but nothing found", nlefttoconn);
32             start_connect(&file[i]);
33             nconn++;
34             nlefttoconn--;
35         }
36
37         rs = rset;
38         ws = wset;
39         n = Select(maxfd + 1, &rs, &ws, NULL, NULL);
40
41         for (i = 0; i < nfiles; i++) {
42             flags = file[i].f_flags;
43             if (flags == 0 || flags & F_DONE)
44                 continue;
45             fd = file[i].f_fd;
46             if (flags & F_CONNECTING &&
47                 (FD_ISSET(fd, &rs) || FD_ISSET(fd, &ws))) {
48                 n = sizeof(error);
```

```
37     ws = wset;
38     n = Select(maxfd + 1, &rs, &ws, NULL, NULL);
39
40     for (i = 0; i < nfiles; i++) {
41         flags = file[i].f_flags;
42         if (flags == 0 || flags & F_DONE)
43             continue;
44         fd = file[i].f_fd;
45         if (flags & F_CONNECTING &&
46             (FD_ISSET(fd, &rs) || FD_ISSET(fd, &ws))) {
47             n = sizeof(error);
48             if (getsockopt(fd, SOL_SOCKET, SO_ERROR, &error, &n) < 0 ||
49                 error != 0) {
50                 err_ret("nonblocking connect failed for %s",
51                         file[i].f_name);
52             }
53             /* connection established */
54             printf("connection established for %s\n", file[i].f_name);
55             FD_CLR(fd, &wset); /* no more writeability test */
56             write_get_cmd(&file[i]); /* write() the GET command */
57         } else if (flags & F_READING && FD_ISSET(fd, &rs)) {
58             if ((n = Read(fd, buf, sizeof(buf))) == 0) {
59                 printf("end-of-file on %s\n", file[i].f_name);
60                 Close(fd);
61                 file[i].f_flags = F_DONE; /* clears F_READING */
62                 FD_CLR(fd, &rset);
63                 nconn--;
64                 nlefttoread--;
65             } else {
66                 printf("read %d bytes from %s\n", n, file[i].f_name);
67             }
68         }
69     }
70 }
```

Using Three Parallel Download Is Good Enough

Figure 16.20. Clock time for various numbers of simultaneous connections.

# simultaneous connections	Clock time (seconds), nonblocking	Clock time (seconds), threads
1	6.0	6.3
2	4.1	4.2
3	3.0	3.1
4	2.8	3.0
5	2.5	2.7
6	2.4	2.5
7	2.3	2.3
8	2.2	2.3
9	2.0	2.2

Blocking Accept with Select() May Cause a Problem

- Suppose that a client initiates a connection and after the connection is set up, the client immediately disconnects it.
 - Why? The server may be too busy to “accept” the connection immediately.
- When the server finally has time to “accept” the connection, the aborted connection has been removed from the kernel, causing the server to block!

```
if (FD_ISSET(listenfd, &rset)) {      /* new client connection */
printf("listening socket readable\n");
sleep(5);
clilen = sizeof(cliaddr);
connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);
```

Figure 16.21 TCP echo client that creates connection and sends an RST.

nonblock/tcpcli03.c

```
1 #include    "unp.h"
2 int
3 main(int argc, char **argv)
4 {
5     int      sockfd;
6     struct linger ling;
7     struct sockaddr_in servaddr;
8
9     if (argc != 2)
10        err_quit("usage: tcpcli <IPaddress>");
11
12     sockfd = Socket(AF_INET, SOCK_STREAM, 0);
13
14     bzero(&servaddr, sizeof(servaddr));
15     servaddr.sin_family = AF_INET;
16     servaddr.sin_port = htons(SERV_PORT);
17     Inet_pton(AF_INET, argv[1], &servaddr.sin_addr);
18
19     Connect(sockfd, (SA *) &servaddr, sizeof(servaddr));
20
21     ling.l_onoff = 1;           /* cause RST to be sent on close() */
22     ling.l_linger = 0;
23     Setsockopt(sockfd, SOL_SOCKET, SO_LINGER, &ling, sizeof(ling));
24     Close(sockfd);
25
26     exit(0);
27 }
```

A Client That Issues RST Immediately After Getting Connected

How to Handle This Problem?

- Always set a listening socket to nonblocking when you use select() to indicate when a connection is ready to be accepted.
- Ignore the following error on the subsequent call to accept(): EWOULDBLOCK (for BSD), etc.