

ioctl Function

- ioctl function
- get_ifi_info function
- ARP cache operations: e.g. Print H/W addresses

ioctl Function

```
#include <unistd.h>
```

```
int ioctl (int fd, int request, ... /* void *arg */ );
```

returns: 0 if OK, -1 on error

type of *arg depends on request

Category	request	Description
socket	SIOCATMARK	at out-of-band mark ?
	SIOCSPGRP	set process ID or process group ID of socket
	SIOCGPGRP	get process ID or process group ID of socket
file	FIONBIO	set/clear nonblocking flag
	FIOASYNC	set/clear asynchronous i/o flag
	FIONREAD	get # bytes in receive buffer
	FIOSETOWN	set process ID or process group ID of file
	FIOGETOWN	get process ID or process group ID of file
ARP	SIOCSARP	create/modify ARP entry
	SIOCGARP	get ARP entry
	SIOCDARP	delete ARP entry

ioctl Function (Cont.)

Category	request	Description
interface	SIOCGIFCONF	get list of all interfaces
	SIOCSIFADDR	set interface address
	SIOCGIFADDR	get interface address
	SIOCSIFFLAGS	set interface flags
	SIOCGIFFLAGS	get interface flags
	SIOCSIFDSTADDR	set point-to-point address
	SIOCGIFDSTADDR	get point-to-point address
	SIOCGIFBRDADDR	get broadcast address
	SIOCSIFBRDADDR	set broadcast address
	SIOCGIFNETMASK	get subnet mask
routing	SIOCSIFNETMASK	set subnet mask
	SIOCGIFMETRIC	get interface metric
	SIOCSIFMETRIC	set interface metric
	SIOC...	(many more; implementation dependent)
	SIOCADDRT	add route
	SIOCDELRT	delete route

get_ifi_info Function

returns a linked list of ifi_info structures

- Create a TCP or UDP socket
- Issue SIOCGIFCONF request in ioctl in a loop
- Initialize linked list pointers
- Step to next socket address structure
- Handle aliases
- Fetch interface flags
- Allocate and initialize ifi_info structure

Figure 17.2 `ifconf` and `ifreq` structures used with various interface `ioctl` requests.

<net/if.h>

```
struct ifconf {
    int ifc_len;                      /* size of buffer, value-result */
    union {
        caddr_t ifcu_buf;            /* input from user -> kernel */
        struct ifreq *ifcu_req;      /* return from kernel -> user */
    } ifc_ifcu;
};

#define ifc_buf ifc_ifcu.ifcu_buf /* buffer address */
#define ifc_req ifc_ifcu.ifcu_req /* array of structures returned */

#define IFNAMSIZ 16

struct ifreq {
    char ifr_name[IFNAMSIZ]; /* interface name, e.g., "le0" */
    union {
        struct sockaddr ifru_addr;
        struct sockaddr ifru_dstaddr;
        struct sockaddr ifru_broadaddr;
        short ifru_flags;
        int ifru_metric;
        caddr_t ifru_data;
    } ifr_ifru;
};

#define ifr_addr ifr_ifru.ifru_addr /* address */
#define ifr_dstaddr ifr_ifru.ifru_dstaddr /* other end of point-to-point */
#define ifr_broadaddr ifr_ifru.ifru_broadaddr /* broadcast address */
#define ifr_flags ifr_ifru.ifru_flags /* flags */
#define ifr_metric ifr_ifru.ifru_metric /* metric */
#define ifr_data ifr_ifru.ifru_data /* for use by interface */
```

Figure 17.3. Initialization of `ifconf` structure before `SIOCGIFCONF`.

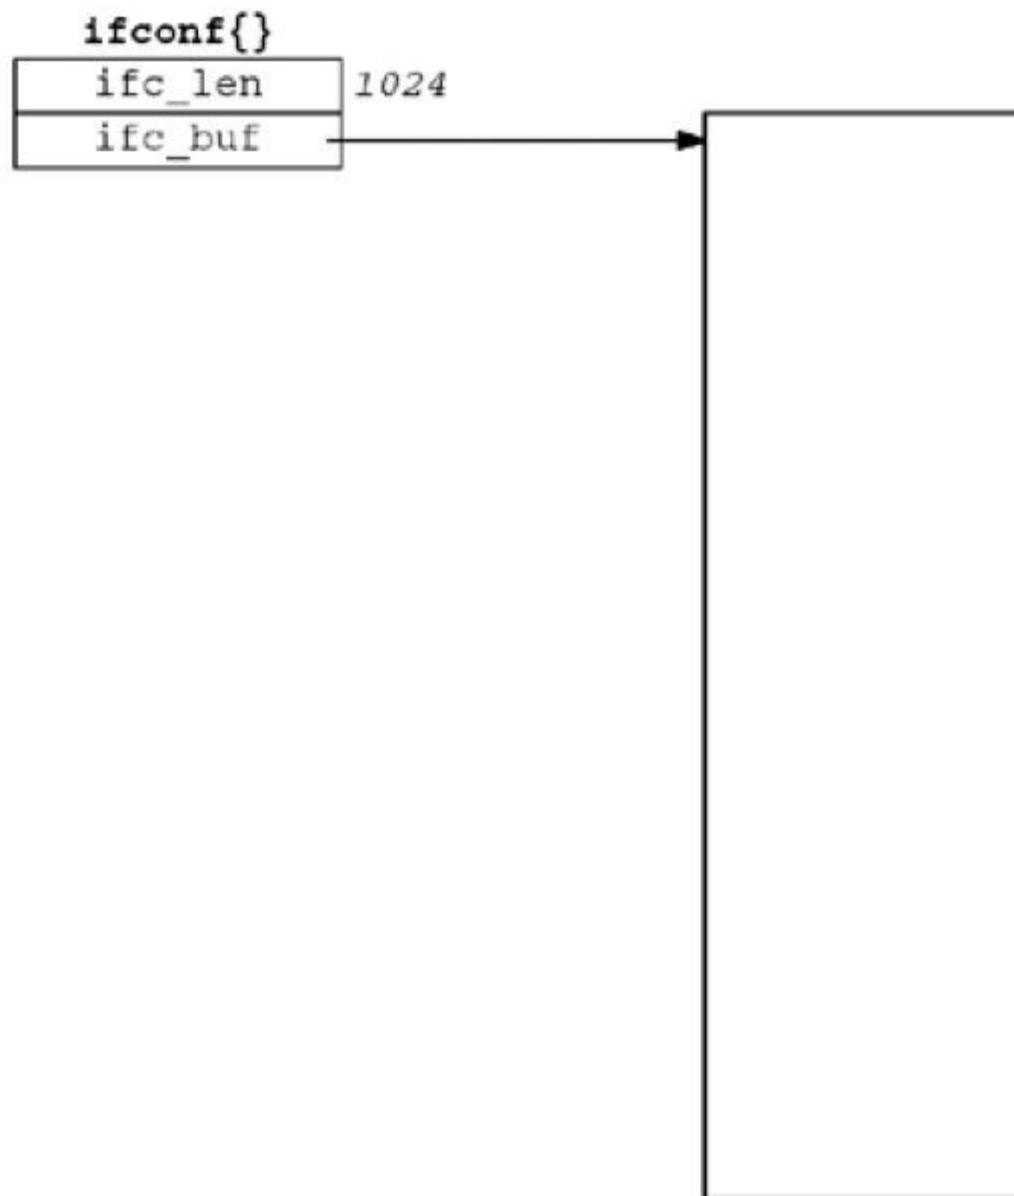
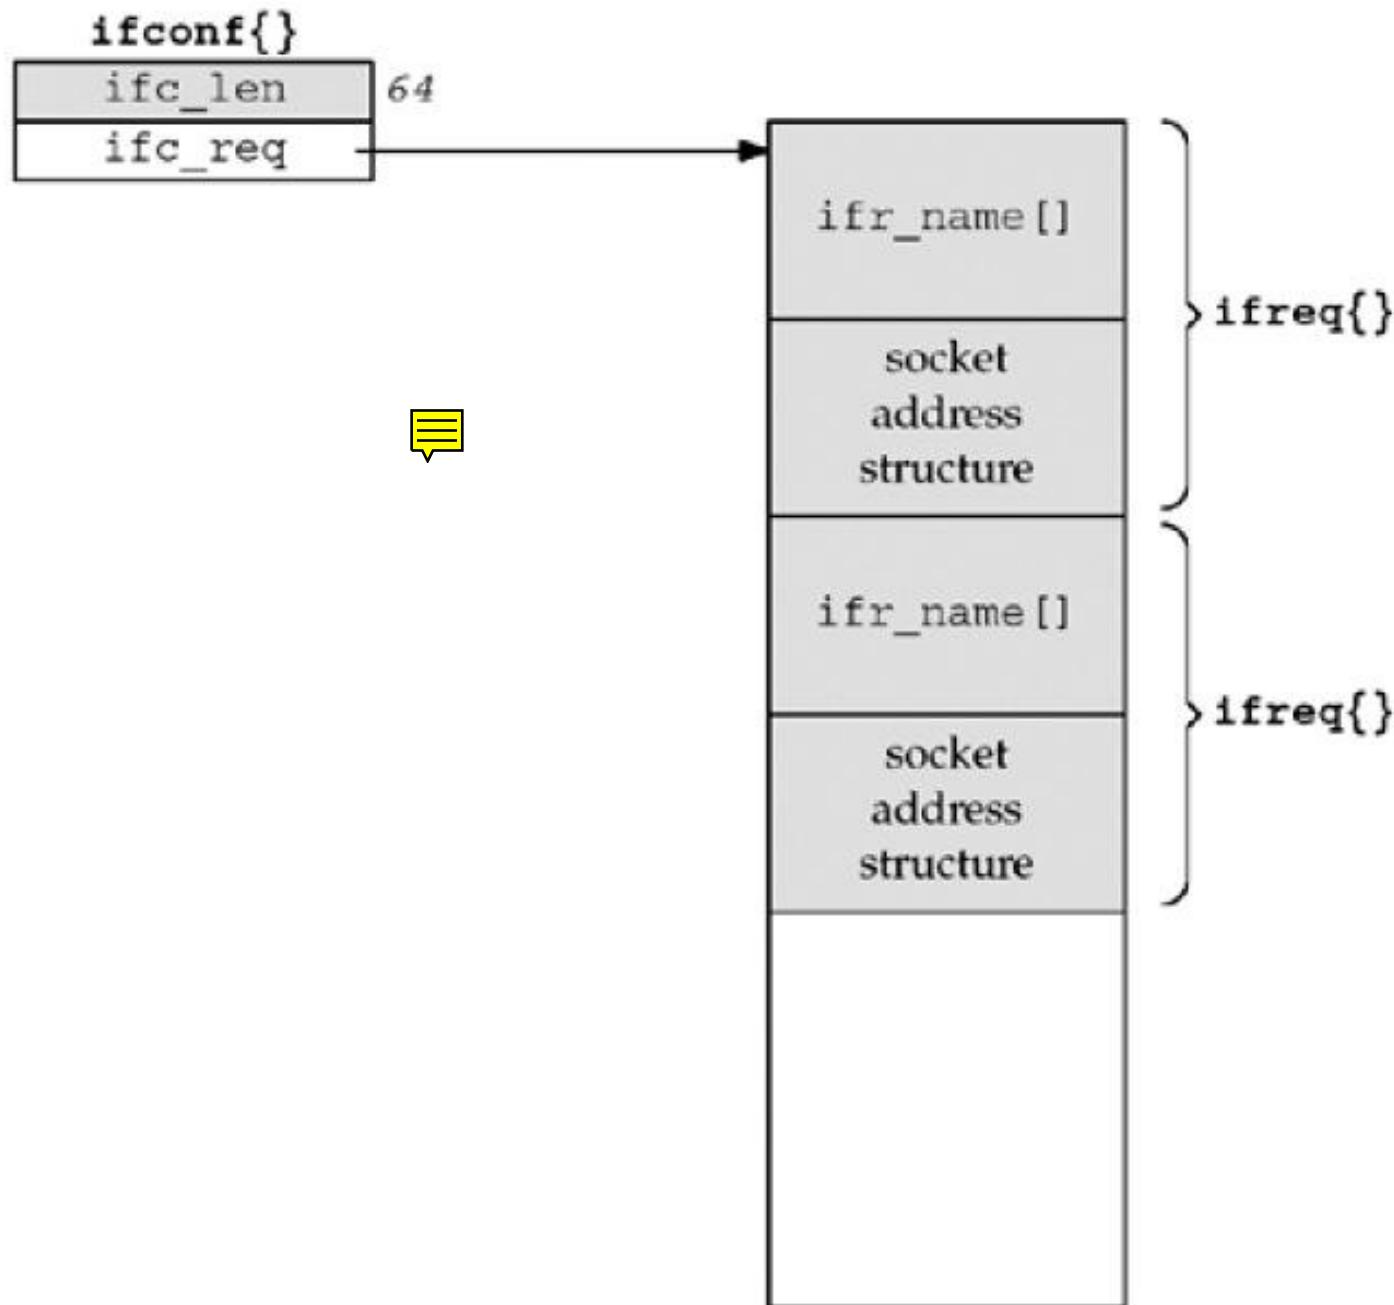


Figure 17.4. Values returned by SIOCGIFCONF.



```
struct sockaddr_dl {
    uint8_t      sdl_len;
    sa_family_t  sdl_family;    /* AF_LINK */
    uint16_t     sdl_index;    /* system assigned index, if > 0 */
    uint8_t      sdl_type;     /* IFT_ETHER, etc. from <net/if_types.h> */
    uint8_t      sdl_nlen;     /* name length, starting in sdl_data[0] */
    uint8_t      sdl_alen;     /* link-layer address length */
    uint8_t      sdl_slen;     /* link-layer selector length */
    char        sdl_data[12];  /* minimum work area, can be larger;
                                contains i/f name and link-layer address */
};
```

Figure 18.1 Datalink socket address structure.



```
macosx % prifinfo inet4 0
lo0: <UP MCAST LOOP >
    MTU: 16384
    IP addr: 127.0.0.1
en1: <UP BCAST MCAST >
    MTU: 1500
    IP addr: 172.24.37.78
    broadcast addr: 172.24.37.95
```

The first command-line argument of `inet4` specifies IPv4 address. The second command-line argument of 0 specifies that no address aliases are to be returned (see return address aliases in Section A.4). Note that under MacOS X, the `l` Ethernet interface is not available using this method.

If we add three alias addresses to the Ethernet interface (`en0`, `en1`, and `en2`, and if we change the second command-line argument to 1, the following:

```
macosx % prifinfo inet4 1
lo0: <UP MCAST LOOP >
    MTU: 16384
    IP addr: 127.0.0.1
en1: <UP BCAST MCAST >
    MTU: 1500
    IP addr: 172.24.37.78          primary IP address
    broadcast addr: 172.24.37.95
en1: <UP BCAST MCAST >
    MTU: 1500
    IP addr: 172.24.37.79          first alias
    broadcast addr: 172.24.37.95
en1: <UP BCAST MCAST >
    MTU: 1500
    IP addr: 172.24.37.80          second alias
    broadcast addr: 172.24.37.95
en1: <UP BCAST MCAST >
    MTU: 1500
    IP addr: 172.24.37.81          third alias
    broadcast addr: 172.24.37.95
```

Figure 17.5 `unpifi.h` header.

`lib/unpifi.h`

```
1 /* Our own header for the programs that need interface configuration info.
2  Include this file, instead of "unp.h". */

3 #ifndef __unp_ifi_h
4 #define __unp_ifi_h

5 #include      "unp.h"
6 #include      <net/if.h>

7 #define IFI_NAME      16          /* same as IFNAMSIZ in <net/if.h> */
8 #define IFI_HADDR     8           /* allow for 64-bit EUI-64 in future */

9 struct ifi_info {
10     char      ifi_name[IFI_NAME];    /* interface name, null-terminated */
11     short     ifi_index;           /* interface index */
12     short     ifi_mtu;            /* interface MTU */
13     u_char    ifi_haddr[IFI_HADDR]; /* hardware address */
14     u_short   ifi_hlen;           /* # bytes in hardware address: 0, 6, 8 */
15     short     ifi_flags;          /* IFF_xxx constants from <net/if.h> */
16     short     ifi_myflags;         /* our own IFI_xxx flags */
17     struct    sockaddr *ifi_addr; /* primary address */
18     struct    sockaddr *ifi_brdaddr; /* broadcast address */
19     struct    sockaddr *ifi_dstaddr; /* destination address */
20     struct    ifi_info *ifi_next;  /* next of these structures */
21 };

22 #define IFI_ALIAS      1           /* ifi_addr is an alias */

23                         /* function prototypes */
24 struct ifi_info *get_ifi_info(int, int);
25 struct ifi_info *Get_ifi_info(int, int);
26 void    free_ifi_info(struct ifi_info *);

27 #endif /* __unp_ifi_h */
```

Figure 17.6 `prifinfo` program that calls our `get_ifi_info` function.

iocctl/prifinfo.c

```
1 #include      "unpifi.h"
2
3 int
4 main(int argc, char **argv)
5 {
6     struct ifi_info *ifi, *ifihead;
7     struct sockaddr *sa;
8     u_char *ptr;
9     int     i, family, doaliases;
10    if (argc != 3)
11        err_quit("usage: prifinfo <inet4|inet6> <doaliases>");
12    if (strcmp(argv[1], "inet4") == 0)
13        family = AF_INET;
14    else if (strcmp (argv[1], "inet6") == 0)
15        family = AF_INET6;
16    else
17        err_quit("invalid <address-family>");
18    doaliases = atoi(argv[2]);
19    for (ifihead = ifi = Get_ifi_info(family, doaliases);
20         ifi != NULL; ifi = ifi->ifi_next) {
21        printf("%s: ", ifi->ifi_name);
22        if (ifi->ifi_index != 0)
23            printf("(%d) ", ifi->ifi_index);
24        printf("<");
25        if (ifi->ifi_flags & IFF_UP)                  printf("UP ");
26        if (ifi->ifi_flags & IFF_BROADCAST)          printf("BCAST ");
27        if (ifi->ifi_flags & IFF_MULTICAST)          printf("MCAST ");
```

```
17 doaliases = atoi(argv[2]);
18 for (ifihead = ifi = Get_ifi_info(family, doaliases);
19      ifi != NULL; ifi = ifi->ifi_next) {
20     printf("%s: ", ifi->ifi_name);
21     if (ifi->ifi_index != 0)
22       printf("(%d) ", ifi->ifi_index);
23     printf("<");
24     if (ifi->ifi_flags & IFF_UP)           printf("UP ");
25     if (ifi->ifi_flags & IFF_BROADCAST)    printf("BCAST ");
26     if (ifi->ifi_flags & IFF_MULTICAST)    printf("MCAST ");
27     if (ifi->ifi_flags & IFF_LOOPBACK)      printf("LOOP ");
28     if (ifi->ifi_flags & IFF_POINTOPOINT)   printf("P2P ");
29     printf(">\n");
30     if ((i = ifi->ifi_hlen) > 0) {
31       ptr = ifi->ifi_haddr;
32       do {
33         printf("%s%x", (i == ifi->ifi_hlen) ? " " : ":", *ptr++);
34       } while (--i > 0);
35       printf("\n");
36     }
37     if (ifi->ifi_mtu != 0)
38       printf(" MTU: %d\n", ifi->ifi_mtu);
39     if ((sa = ifi->ifi_addr) != NULL)
40       printf(" IP addr: %s\n", Sock_ntop_host (sa, sizeof (*sa)));
41     if ((sa = ifi->ifi_brdaddr) != NULL)
42       printf(" broadcast addr: %s\n",
43              Sock_ntop_host (sa, sizeof (*sa)));
44     if ((sa = ifi->ifi_dstaddr) != NULL)
45       printf(" destination addr: %s\n",
46              Sock_ntop_host (sa, sizeof (*sa)));
47   }
48   free_ifi_info(ifihead);
49   exit(0);
```

Figure 17.7 Issue SIOCGIFCONF request to obtain interface configuration.

lib/get_ifi_info.c

```
1 #include      "unpifi.h"
2 struct ifi_info *
3 get_ifi_info(int family, int doaliases)
4 {
5     struct ifi_info *ifi, *ifihead, **ifipnext;
6     int      sockfd, len, lastlen, flags, myflags, idx = 0, hlen = 0;
7     char     *ptr, *buf, lastname[IFNAMSIZ], *cptr, *haddr, *sdlname;
8     struct ifconf ifc;
9     struct ifreq *ifr, ifrcopy;
10    struct sockaddr_in *sinptr;
11    struct sockaddr_in6 *sin6ptr;
12
13    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);
14
15    lastlen = 0;
16    len = 100 * sizeof(struct ifreq); /* initial buffer size guess */
17    for ( ; ; ) {
18        buf = Malloc(len);
19        ifc.ifc_len = len;
20        ifc.ifc_buf = buf;
21        if (ioctl(sockfd, SIOCGIFCONF, &ifc) < 0) {
22            if (errno != EINVAL || lastlen != 0)
23                err_sys("ioctl error");
24        } else {
25            if (ifc.ifc_len == lastlen)
26                break;           /* success, len has not changed */
27            lastlen = ifc.ifc_len;
28        }
29        len += 10 * sizeof(struct ifreq); /* increment */
30        free(buf);
31    }
32    ifihead = NULL;
33    ifipnext = &ifihead;
34    lastname[0] = 0;
```



Figure 17.8 Process interface configuration.

lib/get_ifi_info.c

```
34     for (ptr = buf; ptr < buf + ifc.ifc_len;) {
35         ifr = (struct ifreq *) ptr;
36 #ifdef HAVE_SOCKADDR_SA_LEN
37         len = max(sizeof(struct sockaddr), ifr->ifr_addr.sa_len);
38 #else
39         switch (ifr->ifr_addr.sa_family) {
40 #ifdef IPV6
41             case AF_INET6:
42                 len = sizeof(struct sockaddr_in6);
43                 break;
44 #endif
45             case AF_INET:
46             default:
47                 len = sizeof(struct sockaddr);
48                 break;
49 }
50 #endif /* HAVE_SOCKADDR_SA_LEN */
51         ptr += sizeof(ifr->ifr_name) + len; /* for next one in buffer */
52 #ifdef HAVE_SOCKADDR_DL_STRUCT
53         /* assumes that AF_LINK precedes AF_INET or AF_INET6 */
54         if (ifr->ifr_addr.sa_family == AF_LINK) {
55             struct sockaddr_dl *sdl = (struct sockaddr_dl *) &ifr->ifr_addr;
56             sdlname = ifr->ifr_name;
57             idx = sdl->sdl_index;
58             haddr = sdl->sdl_data + sdl->sdl_nlen;
59             hlen = sdl->sdl_alen;
60         }
61 #endif
62         if (ifr->ifr_addr.sa_family != family)
63             continue; /* ignore if not desired address family */
64 }
```

```
47     len = sizeof(struct sockaddr);
48     break;
49 }
50 #endif /* HAVE_SOCKADDR_SA_LEN */
51     ptr += sizeof(ifr->ifr_name) + len; /* for next one in buffer */

52 #ifdef HAVE_SOCKADDR_DL_STRUCT
53     /* assumes that AF_LINK precedes AF_INET or AF_INET6 */
54     if (ifr->ifr_addr.sa_family == AF_LINK) {
55         struct sockaddr_dl *sdl = (struct sockaddr_dl *) &ifr->ifr_addr;
56         sdlname = ifr->ifr_name;
57         idx = sdl->sdl_index;
58         haddr = sdl->sdl_data + sdl->sdl_nlen;
59         hlen = sdl->sdl_alen;
60     }
61 #endif

62     if (ifr->ifr_addr.sa_family != family)
63         continue; /* ignore if not desired address family */

64     myflags = 0;
65     if ((cptr = strchr(ifr->ifr_name, ':')) != NULL)
66         *cptr = 0; /* replace colon with null */
67     if (strncmp(lastname, ifr->ifr_name, IFNAMSIZ) == 0) {
68         if (doaliases == 0)
69             continue; /* already processed this interface */
70         myflags = IFI_ALIAS;
71     }
72     memcpy(lastname, ifr->ifr_name, IFNAMSIZ);

73     ifrcopy = *ifr;
74     Ioctl(sockfd, SIOCGIFFLAGS, &ifrcopy);
75     flags = ifrcopy.ifr_flags;
76     if ((flags & IFF_UP) == 0)
77         continue; /* ignore if interface not up */
```

Figure 17.9 Allocate and initialize `ifi_info` structure.

`lib/get_ifi_info.c`

```
78         ifi = Calloc(1, sizeof(struct ifi_info));
79         *ifipnext = ifi;           /* prev points to this new one */
80         ifipnext = &ifi->ifi_next; /* pointer to next one goes here */
81
81         ifi->ifi_flags = flags; /* IFF_xxx values */
82         ifi->ifi_myflags = myflags; /* IFI_xxx values */
83 #if defined(SIOCGIFMTU) && defined(HAVE_STRUCT_IFREQ_IFR_MTU)
84         Ioctl(sockfd, SIOCGIFMTU, &ifrcopy);
85         ifi->ifi_mtu = ifrcopy.ifr_mtu;
86 #else
87         ifi->ifi_mtu = 0;
88 #endif
89         memcpy(ifi->ifi_name, ifr->ifr_name, IFI_NAME);
90         ifi->ifi_name [IFI_NAME - 1] = '\0';
91         /* If the sockaddr_dl is from a different interface, ignore it */
92         if (sdlname == NULL || strcmp(sdlname, ifr->ifr_name) != 0)
93             idx = hlen = 0;
94         ifi->ifi_index = idx;
95         ifi->ifi_hlen = hlen;
96         if (ifi->ifi_hlen > IFI_HADDR)
97             ifi->ifi_hlen = IFI_HADDR;
98         if (hlen)
99             memcpy(ifi->ifi_haddr, haddr, ifi->ifi_hlen);
```

Figure 17.10 Fetch and return interface addresses.

lib/get_ifi_info.c

```
100         switch (ifr->ifr_addr.sa_family) {
101             case AF_INET:
102                 sinptr = (struct sockaddr_in *) &ifr->ifr_addr;
103                 ifi->ifi_addr = Calloc(1, sizeof(struct sockaddr_in));
104                 memcpy(ifi->ifi_addr, sinptr, sizeof(struct sockaddr_in));

105 #ifdef SIOCGIFBRDADDR
106             if (flags & IFF_BROADCAST) {
107                 Ioctl(sockfd, SIOCGIFBRDADDR, &ifrcopy);
108                 sinptr = (struct sockaddr_in *) &ifrcopy.ifr_broadaddr;
109                 ifi->ifi_brdaddr = Calloc(1, sizeof(struct sockaddr_in));
110                 memcpy(ifi->ifi_brdaddr, sinptr, sizeof(struct sockaddr_in));
111             }
112 #endif

113 #ifdef SIOCGIFDSTADDR
114             if (flags & IFF_POINTOPOINT) {
115                 Ioctl(sockfd, SIOCGIFDSTADDR, &ifrcopy);
116                 sinptr = (struct sockaddr_in *) &ifrcopy.ifr_dstaddr;
117                 ifi->ifi_dstaddr = Calloc(1, sizeof(struct sockaddr_in));
118                 memcpy(ifi->ifi_dstaddr, sinptr, sizeof(struct sockaddr_in));
119             }
120 #endif
121             break;

122         case AF_INET6:
123             sin6ptr = (struct sockaddr_in6 *) &ifr->ifr_addr;
124             ifi->ifi_addr = Calloc(1, sizeof(struct sockaddr_in6));
125             memcpy(ifi->ifi_addr, sin6ptr, sizeof(struct sockaddr_in6));
```

Figure 17.11 `free_ifi_info` function: frees dynamic memory allocated by `get_ifi_info`.

lib/get_ifi_info.c

```
143 void
144 free_ifi_info(struct ifi_info *ifihead)
145 {
146     struct ifi_info *ifi, *ifinext;
147
148     for (ifi = ifihead; ifi != NULL; ifi = ifinext) {
149         if (ifi->ifi_addr != NULL)
150             free(ifi->ifi_addr);
151         if (ifi->ifi_brdaddr != NULL)
152             free(ifi->ifi_brdaddr);
153         if (ifi->ifi_dstaddr != NULL)
154             free(ifi->ifi_dstaddr);
155         ifinext = ifi->ifi_next; /* can't fetch ifi_next after free() */
156         free(ifi); /* the ifi_info{} itself */
157     }
}
```

ARP Cache Operations

e.g. print hardware addresses of host

- Get list of addresses and loop through each
- Print IP address
- Issue ioctl and print hardware address

```
#include <net/if_arp.h>
struct arpreq {
    struct sockaddr arp_pa; /* protocol address */
    struct sockaddr arp_ha; /* hardware address */
    int     arp_flags; /* flags */
};

#define ATF_INUSE 0x01 /* entry in use */
#define ATF_COM 0x02 /* completed entry (hardware addr valid)*/
#define ATF_PERM 0x04 /* permanent entry */
#define ATF_PUBL 0x08 /* published entry (respond for other host *)
```



Figure 17.12 arpreq structure used with ioctl requests for ARP cache.

<net/if_arp.h>

```
struct arpreq {
    struct sockaddr arp_pa;      /* protocol address */
    struct sockaddr arp_ha;      /* hardware address */
    int             arp_flags;   /* flags */
};

#define ATF_INUSE      0x01 /* entry in use */
#define ATF_COM        0x02 /* completed entry (hardware addr valid) */
#define ATF_PERM       0x04 /* permanent entry */
#define ATF_PUBL       0x08 /* published entry (respond for other host) */
```

Figure 17.13 Print a host's hardware addresses.

ioctl/prmac.c

```
1 #include      "unpifi.h"
2 #include      <net/if_arp.h>
3
4 int
5 main(int argc, char **argv)
6 {
7     int      sockfd;
8     struct ifi_info *ifi;
9     unsigned char *ptr;
10    struct arpreq arpreq;
11    struct sockaddr_in *sin;
12
13    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);
14    for (ifi = get_ifi_info(AF_INET, 0); ifi != NULL; ifi = ifi->ifi_next) {
15        printf("%s: ", Sock_ntop(ifi->ifi_addr, sizeof(struct sockaddr_in)));
16
17        sin = (struct sockaddr_in *) &arpreq.arp_pa;
18        memcpy(sin, ifi->ifi_addr, sizeof(struct sockaddr_in));
19
20        if (ioctl(sockfd, SIOCGARP, &arpreq) < 0) {
21            err_ret("ioctl SIOCGARP");
22            continue;
23        }
24
25        ptr = &arpreq.arp_ha.sa_data[0];
26        printf("%x:%x:%x:%x:%x\n", *ptr, *(ptr + 1),
27               *(ptr + 2), *(ptr + 3), *(ptr + 4), *(ptr + 5));
28    }
29    exit(0);
30 }
```

```
hpx % prmac
192.6.38.100: 0:60:b0:c2:68:9b
192.168.1.1: 0:60:b0:b2:28:2b
127.0.0.1: ioctl SIOCGARP: Invalid argument
```